

Beyond Security: Leveraging OPA for FinOps in Kubernetes



Open Policy Agent

SATHISH KUMAR VENKATESAN

#whoami

SATHISH KUMAR VENKATESAN

DevOpsCloudJunction - Founder & CEO



LinkedIn



Website

Interests and Hobbies:

- Passionate about computers, gadgets, electronics, photography, cloud computing, DevOps, and open-source technologies.
- Enjoy participating in tech communities and contributing to collaborative projects.

What I'll Cover Today

The FinOps challenge in Kubernetes environments

OPA fundamentals and capabilities

Extending OPA beyond security to cost governance

Practical FinOps policies with Rego

Adopting strategies and best practices

Live demo: OPA FinOps in action

FinOps

FinOps is the practice of bringing financial accountability to the cloud's dynamic environments.

Key Goals:

- Inform: Create visibility into cloud costs and usage.
- Optimize: Maximize cloud value through right-sizing and reservations
- Operate: Establish continuous processes for financial control



The FinOps Challenge

Average Kubernetes cloud waste: 32% of resources underutilized

Common issues:

- Overprovisioned resources
- Idle workloads
- Non-optimal node selection
- Missing resource quotas
- Untagged resources

CPU and memory utilization


The average CPU utilization across clusters remained low at 10% (-23% YoY), while average memory utilization was marginally better at 23% (+15% YoY), indicating no significant year-over-year improvement in resource efficiency across cloud platforms compared to our previous report from 2024.



Cloud provider	aws		
CPU utilization (resources used vs. provisioned)	10%	12%	8%
Memory utilization (resources used vs. provisioned)	24%	23%	21%
CPU overprovisioning (resources requested vs. provisioned)	39%	39%	41%
Memory overprovisioning (resources requested vs. provisioned)	58%	53%	65%

The Cost of Kubernetes Inefficiency

- Average CPU utilization remains critically low at 10%
- Memory utilization hovers at 23% across cloud providers
- Over 50% of requested memory resources remain unused



What is Open Policy Agent (OPA)?

What are Policies in OPA?

```
package example

allow {
  input.user == "alice"
  input.method == "GET"
  input.path[0] == "users"
}
```

Rules in Rego: Policies are rules written in the Rego language.

Evaluate Data: Policies make decisions based on input data

Allow or Deny: The core outcome is typically to allow or deny an action

Policy as Code: Policies are text, enabling version control

Decoupled Logic: Policy decisions are separate from applications

Centralized Control: OPA provides a single place to manage policies

Why OPA for FinOps?



Reuse existing security tooling



Single policy framework for multiple domains



Shift-left approach to cost management



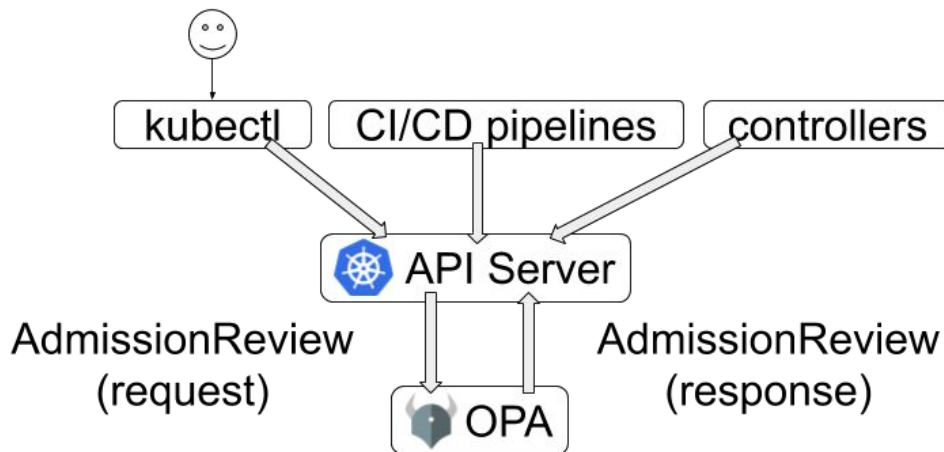
Consistent enforcement across environments





GitOps-friendly policies as code

What is OPA Gatekeeper?

- Kubernetes-native implementation of OPA that enforces policies on Kubernetes resources
- Uses CustomResourceDefinitions (CRDs) to define and manage policies within the Kubernetes ecosystem
- Validates resources against policies during admission control before they're created or modified
- Provides audit capabilities to detect and report on existing policy violations in a cluster



Traditional Vs FinOps - OPA Use Cases



```
package templates.resourcerequests

violation[{"msg": msg}] {
  container := input.review.object.spec.containers[_]
  cpu_request := container.resources.requests.cpu
  memory_request := container.resources.requests.memory

  # Convert CPU to millicores for calculation
  cpu_millicore := to_number(cpu_request) * 1000

  # Check if CPU request exceeds allowed thresholds
  cpu_millicore > data.parameters.maxCPUMillicores

  msg := sprintf("Container '%v' CPU request of '%v' exceeds maximum allowed (%v millicores)",
    [container.name, cpu_request, data.parameters.maxCPUMillicores])
}
```

Example: Resource Request Optimization

GOAL: PREVENT
OVER-PROVISIONING OF
RESOURCES

Example: Enforce Spot Instances for Dev Workloads

GOAL: ENSURE ALL DEVELOPMENT
WORKLOADS RUN ON COST-EFFICIENT
SPOT INSTANCES

```
package kubernetes.finops.spot_instances

violation[{"msg": msg}] {
    input.review.kind.kind == "Pod"

    # Identify dev workloads via namespace or labels
    is_dev_workload

    # Check if spot instance selector is missing
    not has_spot_instance_selector

    msg := sprintf("Dev workload '%v' must use spot instances",
        [input.review.object.metadata.name])
}

# Match dev environments
is_dev_workload {
    contains(input.review.object.metadata.namespace, "dev")
}

is_dev_workload {
    input.review.object.metadata.labels.environment == "dev"
}

# Check for appropriate spot instance configuration
has_spot_instance_selector {
    input.review.object.spec.nodeSelector["node.kubernetes.io/lifecycle"] == "spot"
}
```

Example: Mandatory Cost Allocation Tagging Policy

GOAL: ENSURE CONSISTENT LABELING
FOR ACCURATE COST TRACKING,
ATTRIBUTION, AND FINOPS REPORTING

```
package finops.tagging

# Required cost labels for all resources
required_labels = {
    "cost-center",
    "team",
    "environment",
    "project-id"
}

# Check for missing labels on namespaces
violation[{"msg": msg}] {
    input.request.kind.kind == "Namespace"
    missing := required_labels - {label | label := object.get(input.request.object.metadata.labels, [_],
    "")}
    count(missing) > 0
    msg := sprintf("Namespace missing required labels: %v", [concat(", ", missing)])
}

# Check for missing labels on workload resources
violation[{"msg": msg}] {
    workload_kinds := {"Deployment", "StatefulSet", "DaemonSet", "Job", "CronJob"}
    input.request.kind.kind == workload_kinds[_]
    missing := required_labels - {label | label := object.get(input.request.object.metadata.labels, [_],
    "")}
    count(missing) > 0
    msg := sprintf("%v missing required labels: %v", [input.request.kind.kind, concat(", ", missing)])
}
```

```
package budgets

violation[{"msg": msg}] {
  # Get team from namespace labels
  team := input.review.object.metadata.namespace.labels.team

  # Calculate estimated monthly cost for this workload
  estimated_cost := calculate_workload_cost(input.review.object)

  # Compare against team's remaining budget from external data
  team_budget := data.external.budgets[team].remaining
  estimated_cost > team_budget

  msg := sprintf("Deployment would exceed team '%v' remaining budget of $%v", [team, team_budget])
}
```

Example: Budget Enforcement

GOAL: ENFORCE TEAM-LEVEL
SPENDING LIMITS

Example: Storage Optimization

GOAL: PREVENT USE OF PREMIUM
STORAGE FOR NON-CRITICAL DATA

```
package storage

violation[{"msg": msg}] {
    # Check if it's a PVC request
    input.review.kind.kind == "PersistentVolumeClaim"

    # Get storage class
    storage_class := input.review.object.spec.storageClassName

    # Check if premium SSD
    premium_ssd_classes := {"premium-ssd", "io1", "premium-lrs", "ultra-disk"}
    storage_class == premium_ssd_classes[_]

    # Check for justification
    not input.review.object.metadata.annotations["finops.company.com/premium-storage-justification"]

    # Not in production critical tier
    not (input.review.object.metadata.namespace == "production" and
        input.review.object.metadata.labels.tier == "critical")

    msg := sprintf("Premium SSD storage class '%v' requires justification", [storage_class])
}
```


Integration with External Cost Tools

Example: OpenCost Integration

- Connect real-time cost metrics to OPA decisions

Example: Pulling cost data from OpenCost API

```
# External data connector for OpenCost
package finops.opencost

import future.keywords.in

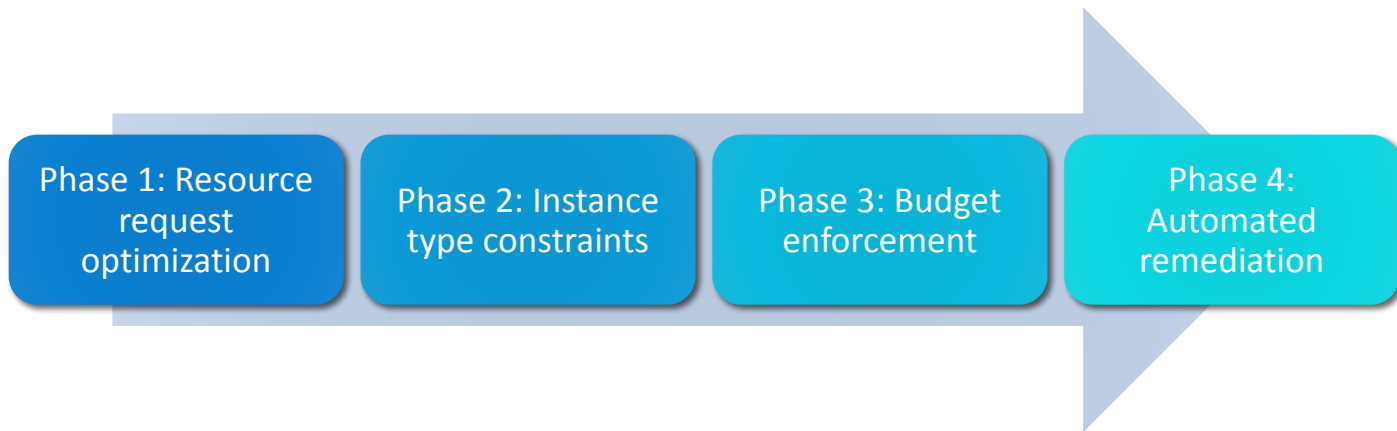
# Calculate estimated cost based on resources
estimated_cost := cost {
  # Get cost data from OpenCost
  cost_data := http.send({
    "method": "GET",
    "url": concat("", ["http://opencost.monitoring.svc:9003/allocation/compute",
                      "?window=1d&aggregate=namespace"]),
    "timeout": "1s"
  }).body

  # Extract cost for requested namespace
  namespace := input.review.object.metadata.namespace
  cost := cost_data.namespaces[namespace].totalCost
}
```

Start Small:

- Begin with visibility-only policies (dry-run mode)
- Focus on highest-cost resources first (GPU, high-memory)

Phased Implementation:



Adoption Strategy & Best Practices

Adoption Strategy & Best Practices



Team Collaboration:

Involve platform, security, finance, and application teams

Create exception processes

Document and communicate policy decisions



Measuring Success:

Track cost reduction metrics

Monitor policy violations and trends

Calculate ROI of policy enforcement



Common Challenges & Solutions

Resources & Next Steps



GitHub Repository:

[https://github.com/open-policy-agent/
opa
gatekeeper-library](https://github.com/open-policy-agent/gatekeeper-library)



Documentation:

OPA/Gatekeeper:

<https://www.openpolicyagent.org/>
[https://open-policy-agent.github.io/gat
ekeeper/website/](https://open-policy-agent.github.io/gatekeeper/website/)



Community:

[OPA Slack:
https://slack.openpolicyagent.org/](https://slack.openpolicyagent.org/)



Demo

Q & A



Feedback

