Explain How Kubernetes Works with GPUs Like I’m 5

Carlos Santana

Sr. WW Spec. SA, Containers at AWS
CNCF Ambassador

Learning at home

kubernetes

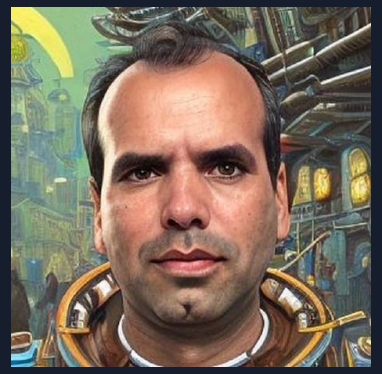

k8s
experience

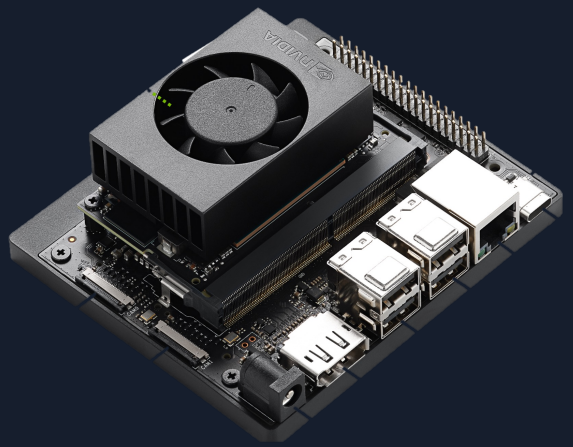

NVIDIA
Jetson GPU

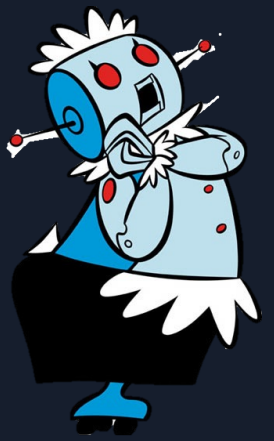

Generative
AI & ML

Images:
https://warnerbros.fandom.com/wiki/Rosey
https://developer.nvidia.com/embedded/learn/get-started-jetson-orin-nano-devkit

Learn and Be Curious

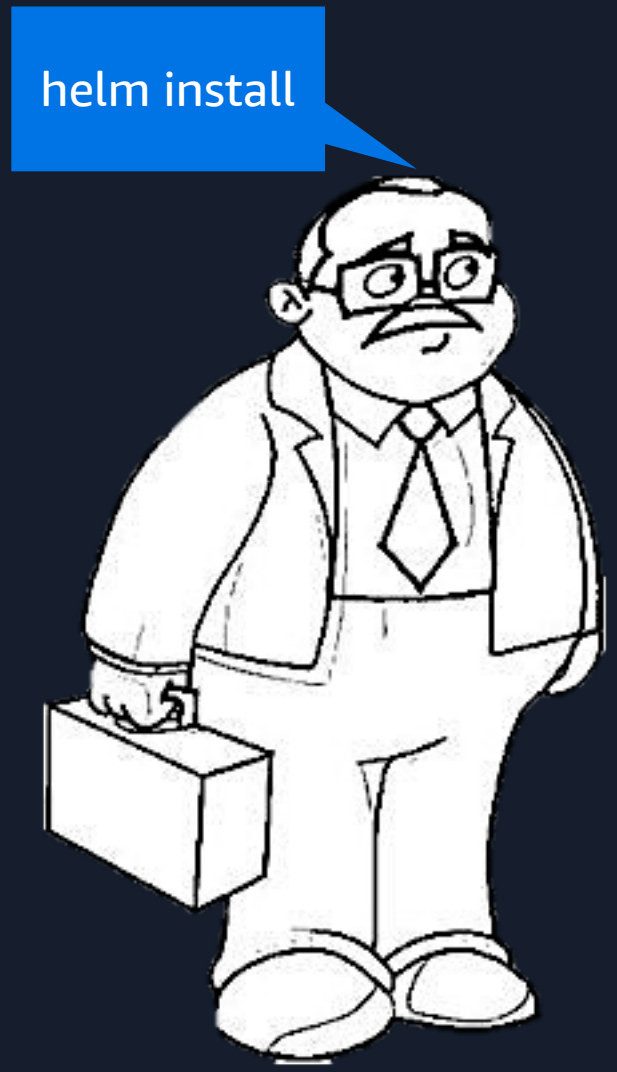

- CUDA
- Container Toolkit
- Device Driver
- Device Plugin
- GPU Feature Discovery
- Node Feature Discovery

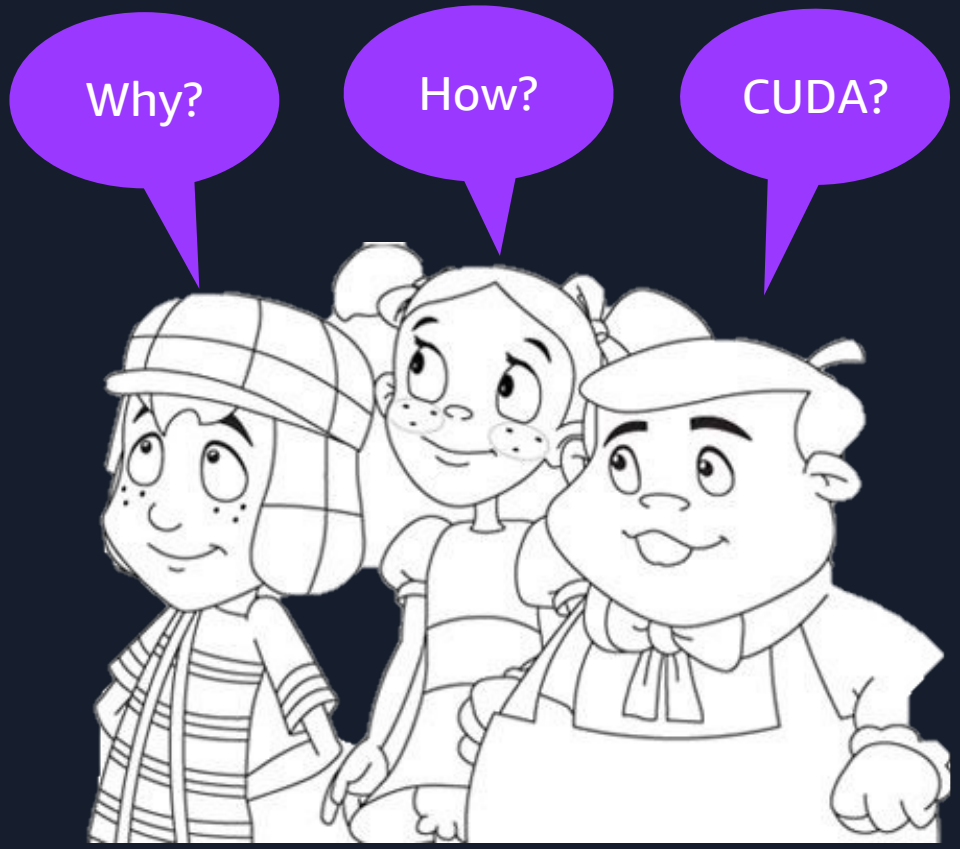

Images:
https://www.pinterest.com/pin/879116789737240075/
https://www.pinterest.com/pin/433401164162584656/

Main areas

HOST
Operating System
Containerd

- Device Driver
- Container Toolkit

CONTAINER
Programing model

- CUDA

KUBERNETES
Kubelet
Node Labels

- Device Plugin
- Node Feature Discovery
- GPU Feature Discovery

Device Driver

- Implements the CUDA Driver API
- Use Vendor pre-installed drivers
- NVIDIA Jetson JetPack OS
- Cloud Providers include driver in OS images
  - Example:
    - AWS EKS AIM EKS AL23, BR, Auto Mode
- Driver package include CUDA user mode driver
- Do not use GPU Operator unless you know why

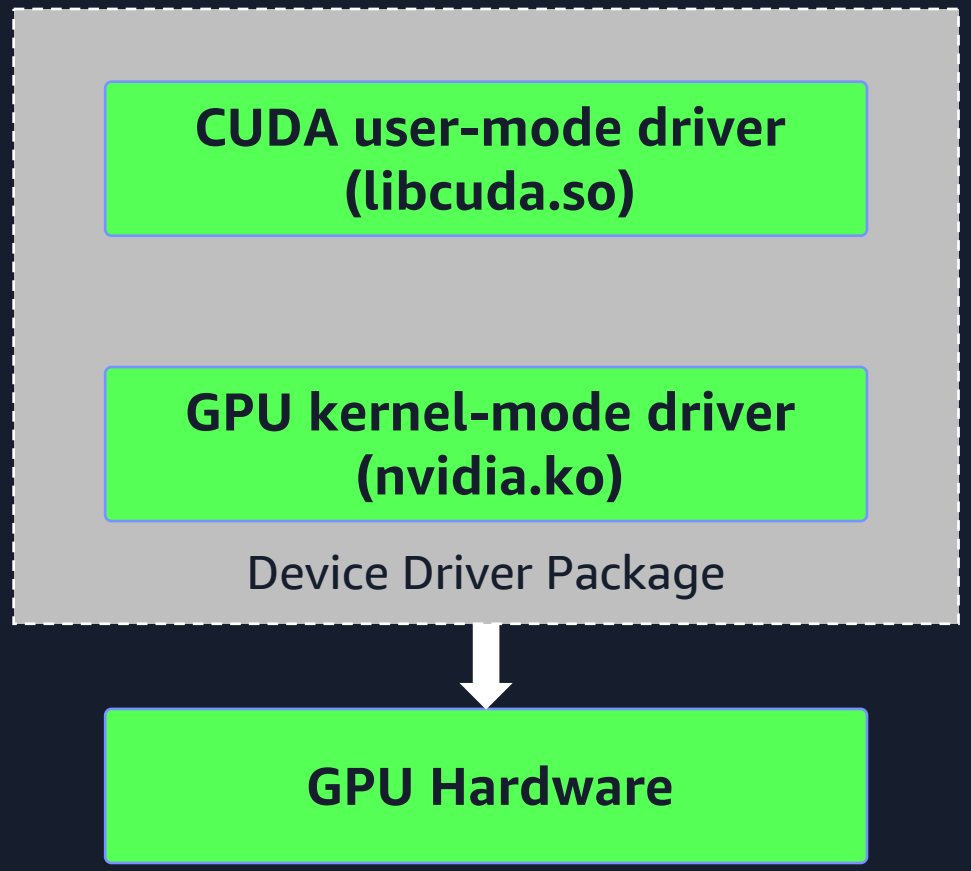

CUDA Toolkit

- Toolkit at build time necessary for linking
- Compatibility between CUDA and Driver
- Container Image usually comes with CUDA runtime
- Do not mount toolkit runtime from host path, unless you know why

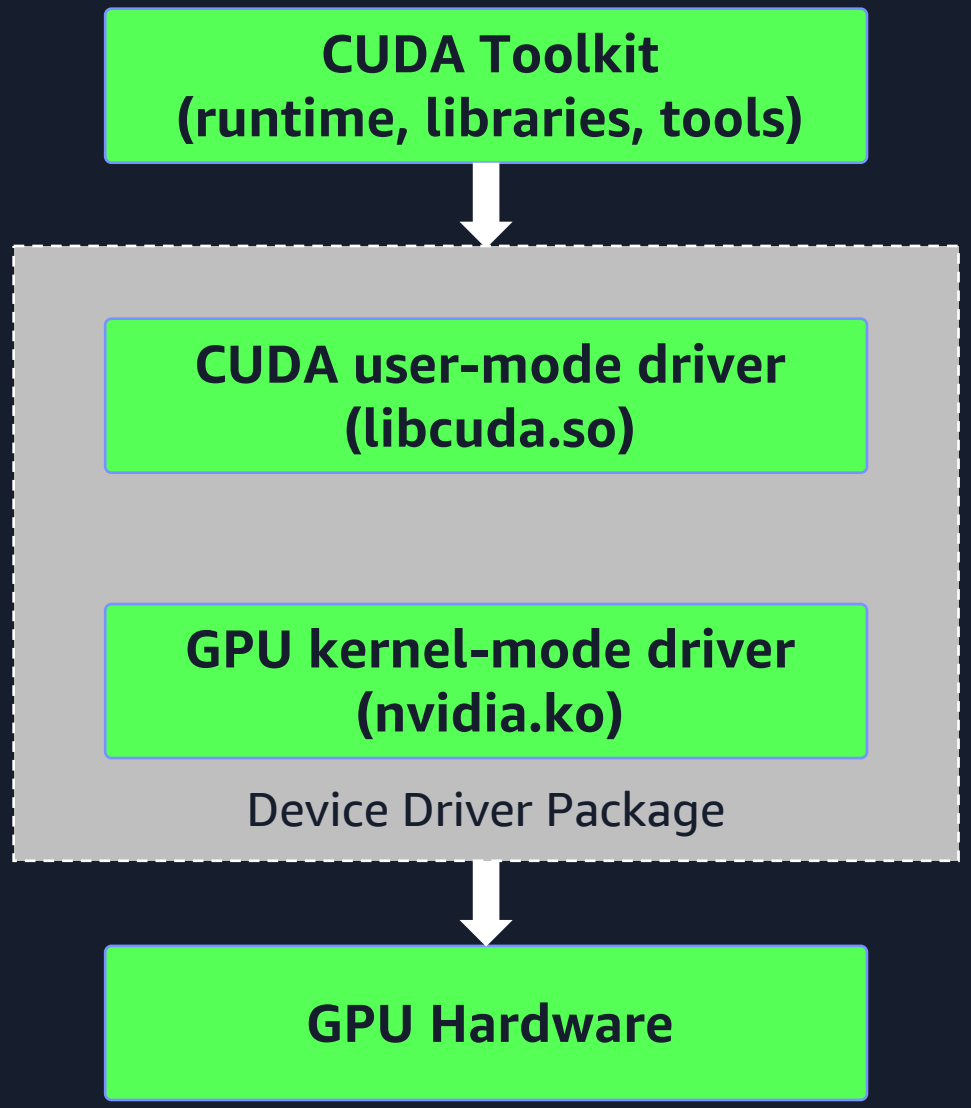

Driver and CUDA versions

nvidia-smi for PCI GPUs
On Jetson dGPU:
/proc/driver/nvidia/version

```
$ cat /proc/driver/nvidia/version
VRM version: NVIDIA UNIX Open Kernel
Module for aarch64  540.3.0

$ head -n 1 /etc/nv_tegra_release
# R36 (release), REVISION: 3.0 (aka 36.3)
```

nvcc (CUDA compiler)
ldconfig

```
$ nvcc --version
Cuda compilation tools, release 12.2, V12.2.140

$ ldconfig –p | grep libcudart
libcudart.so.12 (libc6,AArch64)
```

CUDA Toolkit	Minimum Required Driver Version
CUDA 12.x	>=525.60.13
CUDA 11.x	>=450.80.02

Container Toolkit

- Run GPU-accelerated containers
- Cloud Providers includes this
- For Jetson you need to install docker engine (containerd) and install the toolkit

```
$ nvidia-ctk runtime configure  --runtime=containerd
$ cat /etc/containerd/config.toml
default_runtime_name = "nvidia"
[plugins."io.containerd.grpc.v1.cri".containerd.runtimes.nvidia.options]
  BinaryName = "/usr/bin/nvidia-container-runtime"
```

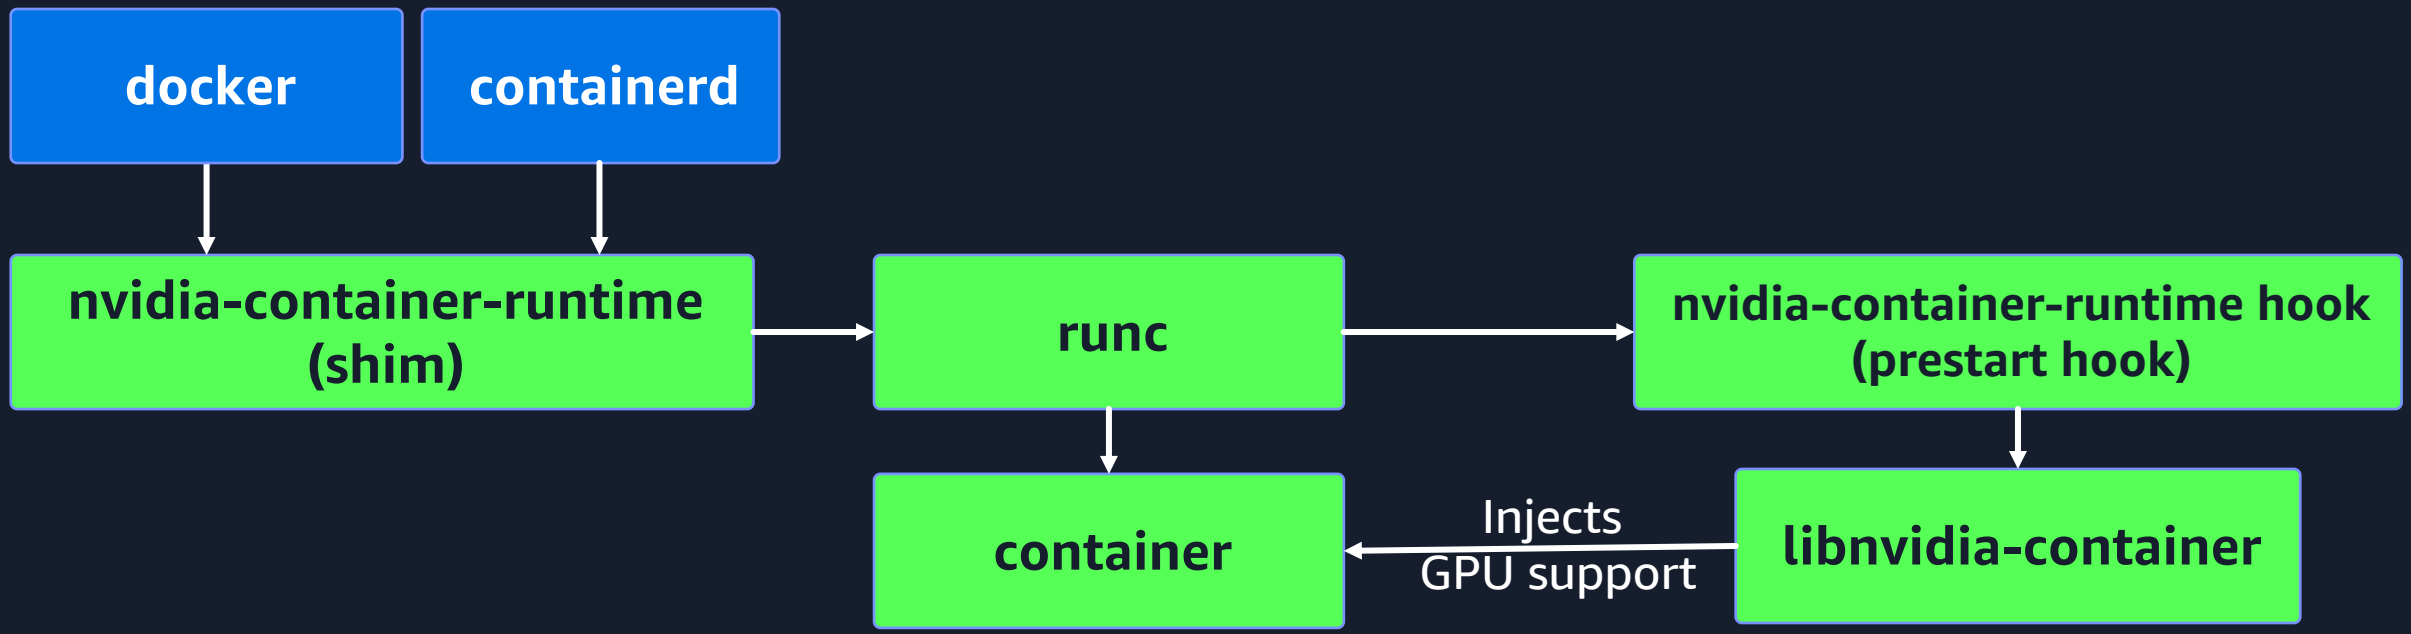

GPU Containers

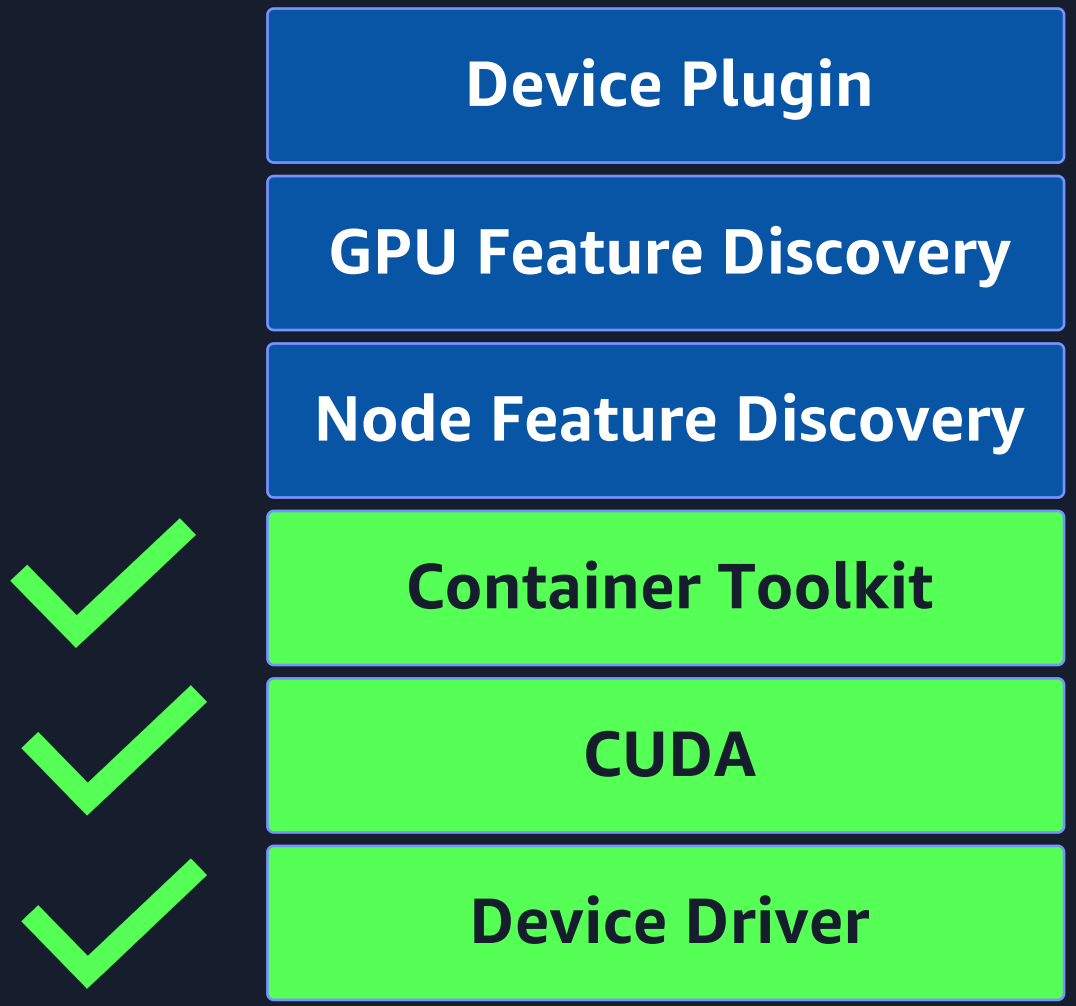


```
$ docker run --runtime nvidia --gpus all
```

Jetson Containers

https://github.com/dusty-nv/jetson-containers

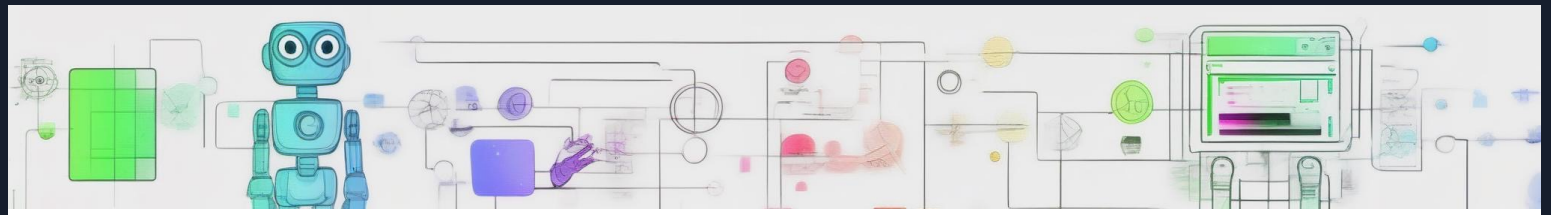

ML	pytorch tensorflow jax onnxruntime deepstream holoscan CTranslate2 JupyterLab
LLM	SGLang vLLM MLC AWQ transformers text-generation-webui ollama llama.cpp llama-factory exllama AutoGPTQ FlashAttention DeepSpeed bitsandbytes xformers
VLM	llava llama-vision VILA LITA NanoLLM ShapeLLM Prismatic xtuner
VIT	NanoOWL NanoSAM Segment Anything (SAM) Track Anything (TAM) clip_trt
RAG	llama-index langchain jetson-copilot NanoDB FAISS RAFT
L4T	l4t-pytorch l4t-tensorflow l4t-ml l4t-diffusion l4t-text-generation
CUDA	cupy cuda-python pycuda numba opencv:cuda cudf cuml
Robotics	Cosmos Genesis ROS LeRobot OpenVLA 3D Diffusion Policy Crossformer MimicGen OpenDroneMap ZED
Graphics	stable-diffusion-webui comfyui nerfstudio meshlab pixsfm gsplat
Mamba	mamba mambavision cobra dimba videomambasuite
Speech	whisper whisper_trt piper riva audiocraft voicecraft xtts
Home/IoT	homeassistant-core wyoming-whisper wyoming-openwakeword wyoming-piper

Jetson Containers

- Modular container build system
- Latest AI/ML packages
- Wraps docker run --runtime=nvidia
- Discord and community meetings

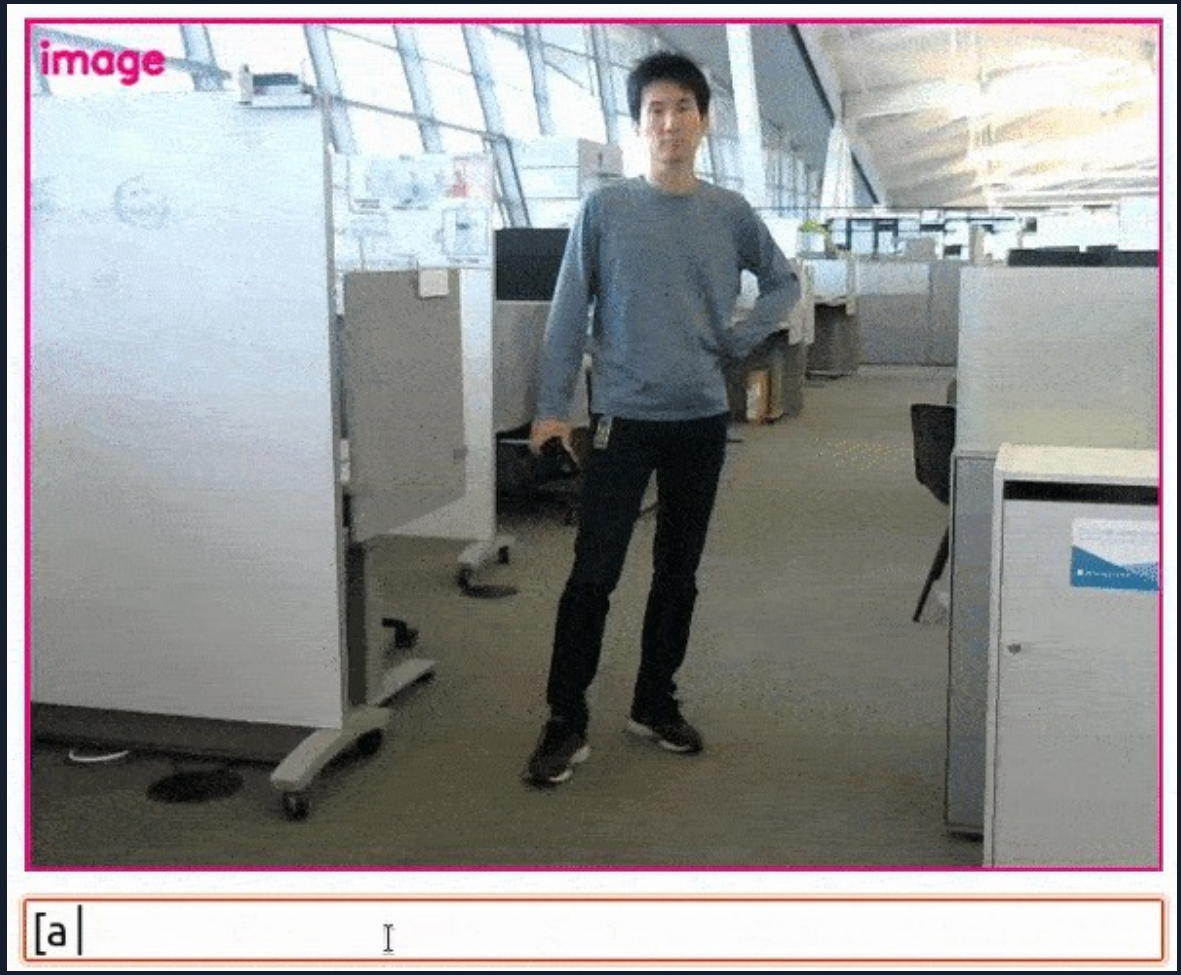

Jetson Containers

$ jetson-containers run --name ollama $(autotag ollama)

>>> How many kubernetes administrators takes to deploy a cluster?

In a typical scenario, a small team of 2-3 administrators may be sufficient for deploying and managing a Kubernetes cluster in a simple environment.
However, larger organizations or complex deployments could require teams of 5-10 or more administrators to ensure proper management and scalability.

Kubernetes

- Device Plugin required?
  - No scheduling
  - There is no allocations
  - All Pods access all GPUs

✅ **Container Toolkit**

✅ **CUDA**

✅ **Device Driver**

```
$ kubectl run myllm --image myllmImage
```

Without allocatable nvidia.com/gpu

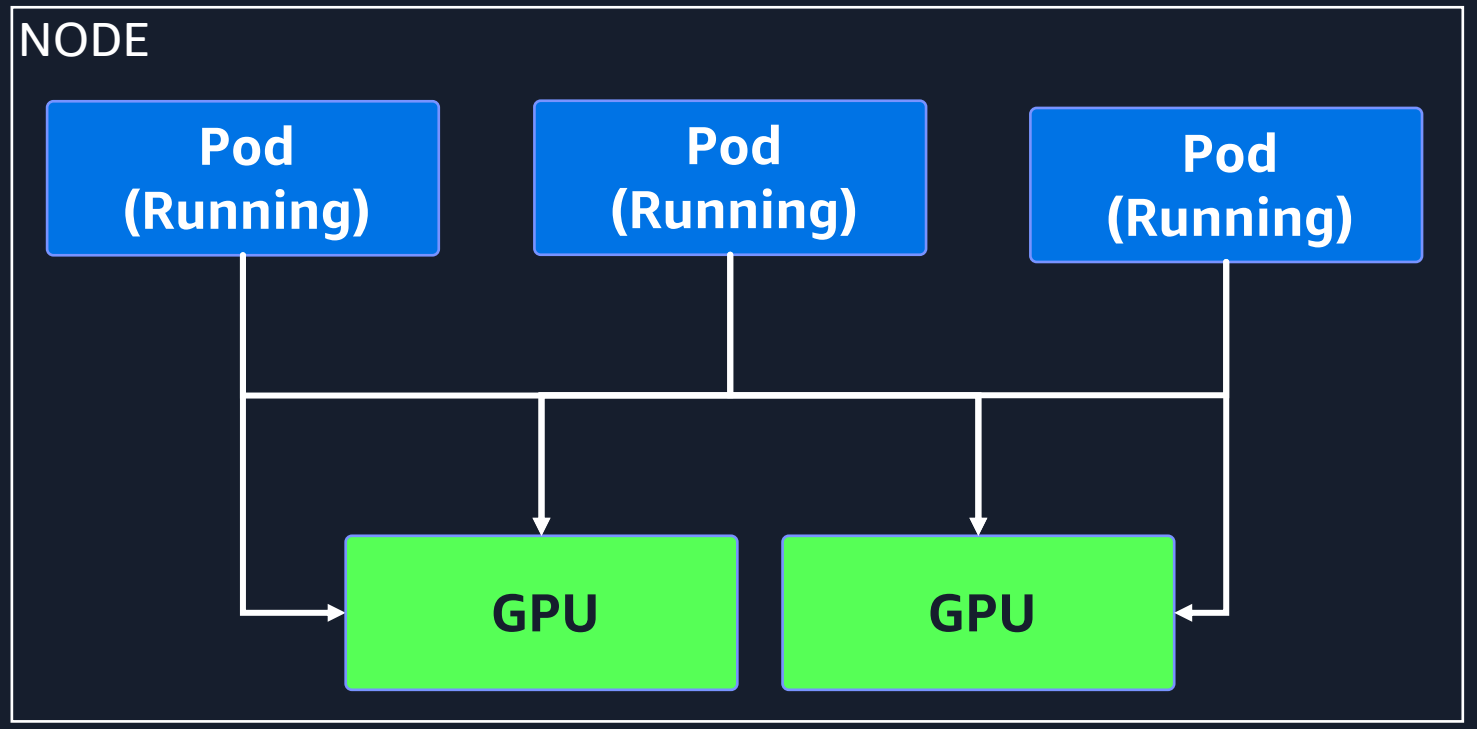

Kubernetes + GPU Device Plugin

- NVIDIA Device Plugin
  - Helm Chart
- Schedule Pods
- Allocate GPUs
- Share a GPU
  - Time-slicing
  - MPS
  - MIG

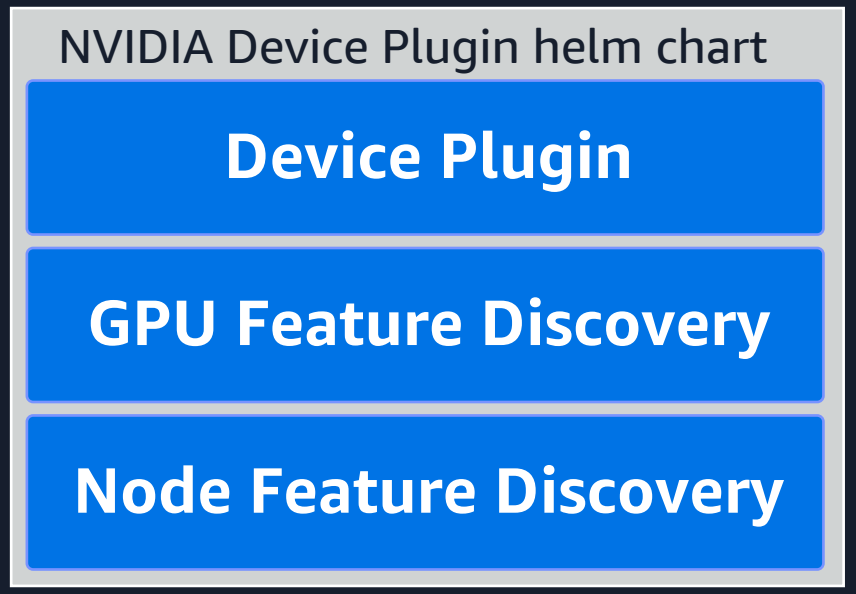

Node Feature Discovery (NFD)

- Advertise features using labels
- Detects hardware on node
- CPU, Kernel, OS, PCI
- PCI VENDOR NVIDIA (10ed)

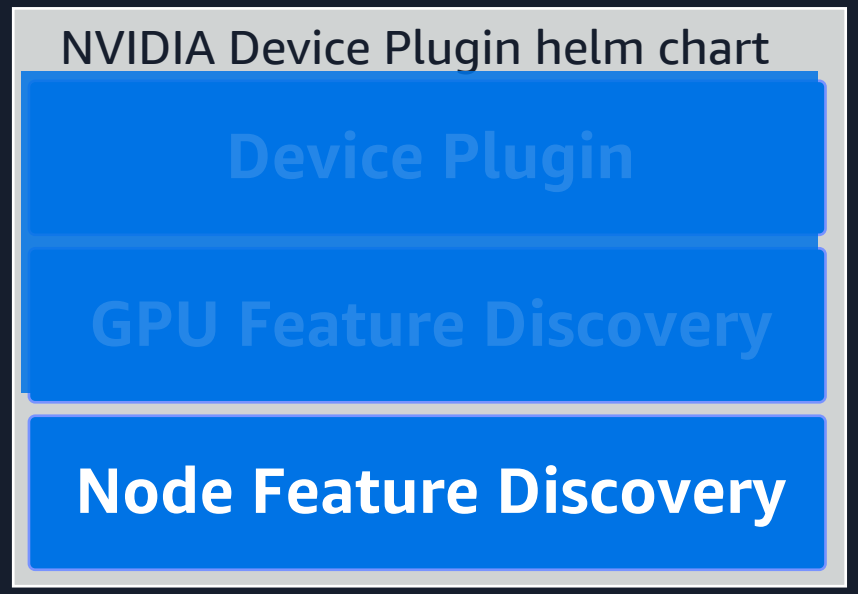

Enables the GPU Feature Discovery →

```
$ kubectl get no jetson -o json

 "cpu-cpuid.USCAT": "true",
 "cpu-model.vendor_id": "ARM",
 "kernel-config.PREEMPT": "true",
 "kernel-version.full": "5.15.136-tegra",
 "storage-nonrotationaldisk": "true",
 "system-os_release.ID": "ubuntu",
 "system-os_release.VERSION_ID": "22.04",

 "pci-10ed.present": "true",
```

GPU Feature Discovery (GFD)

- Advertise features using labels
- Detects NVIDIA GPU
- Labels with nvidia.com/*

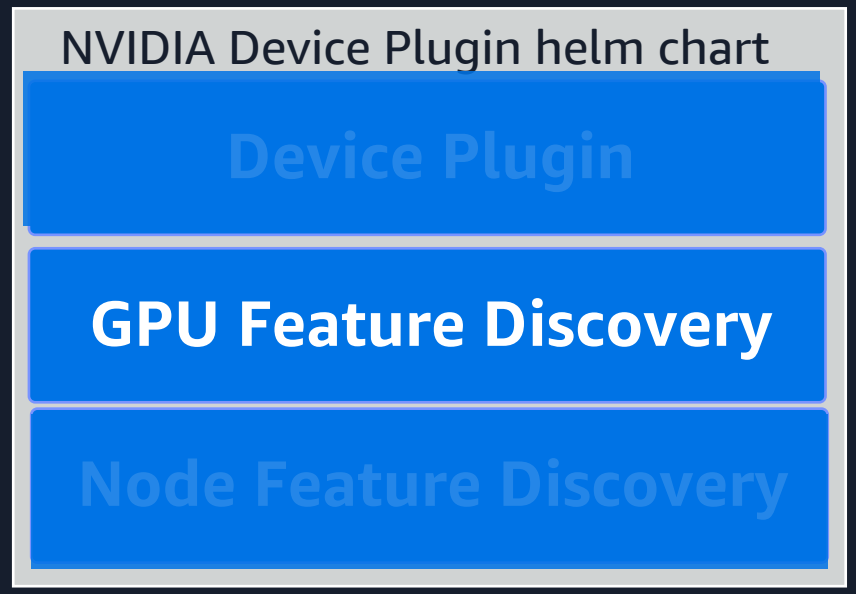

Enables the Device Plugin deamon →

```
$ kubectl get no jetson -o json

"nvidia.com/cuda.driver-version.full": ”540.3.0",
"nvidia.com/cuda.runtime-version.full": "12.2",
"nvidia.com/gpu.count": "1",
"nvidia.com/gpu.replicas": "4",
"nvidia.com/gpu.product": "Orin-SHARED",
"nvidia.com/gpu.sharing-strategy": "time-slicing",
"nvidia.com/mig.capable": "false",
"nvidia.com/mps.capable": "false",
"nvidia.com/vgpu.present": "false"

"nvidia.com/gpu.present": "true",
```

Device Plugin

- Kubelet Device Plugin API
- Expose the number of GPUs
- Tracks health of the GPUs

Kubelet patches this status

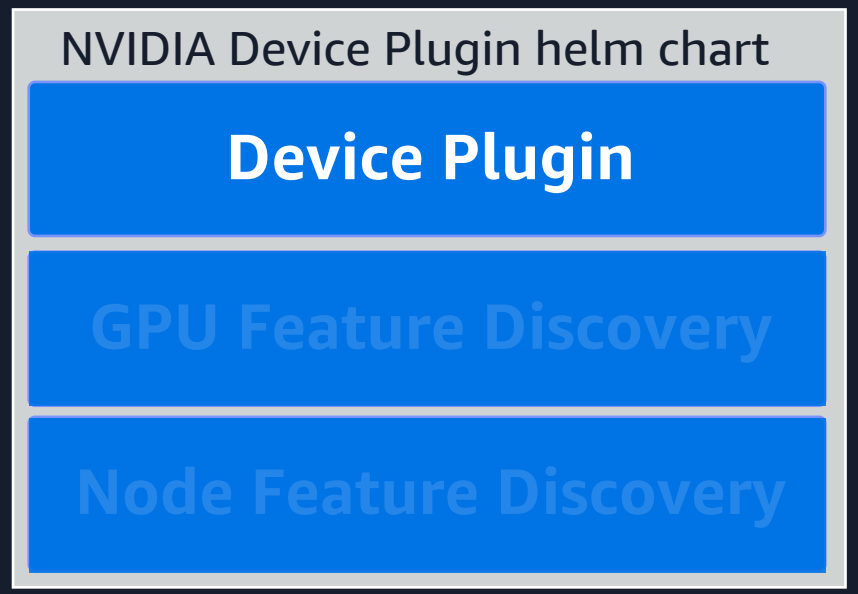


```
k get nodes jetson -o json | jq .status.allocatable
{
  "cpu": "8",
  ”memory": "16032384Ki",
  "nvidia.com/gpu": "4",
  "pods": "110"
}
```

`limit` instead of `request`

```
containers:
- name: cuda-container
  image: nvcr.io/nvidia/k8s/cuda-sample
  resources:
    limits:
      nvidia.com/gpu: 1 # requesting 1 GPU
```

Device Plugin

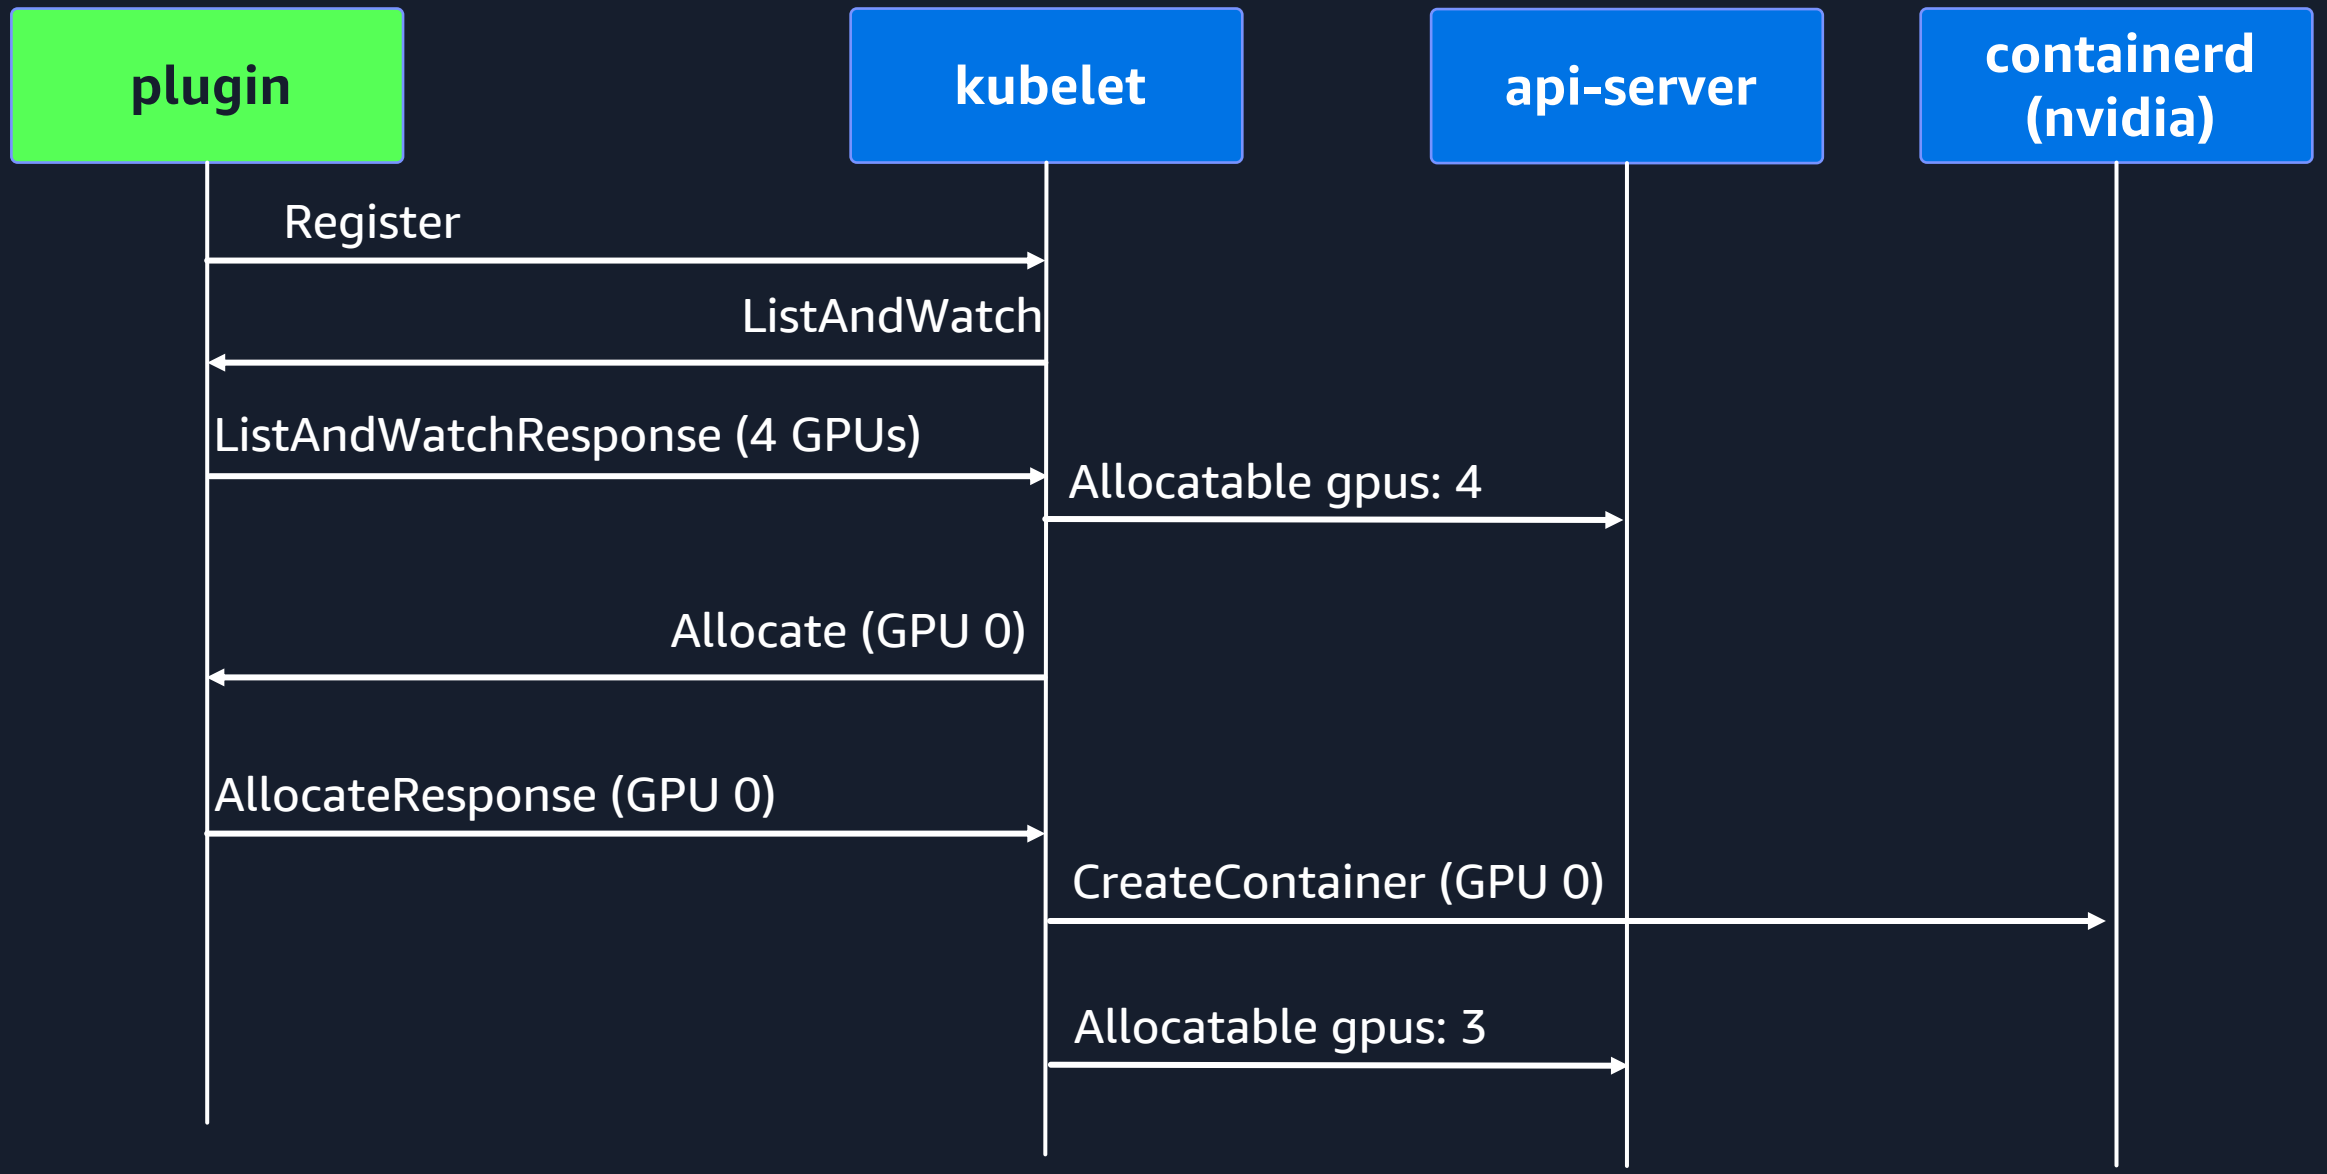

Scheduling Pods

- Allocatable status on Node
- Pods states how much GPU they need
- As Pod wait for GPU, they are pending

```
containers:
- name: cuda-container
  resources:
    limits:
      nvidia.com/gpu: 1 # requesting 1 GPU
```

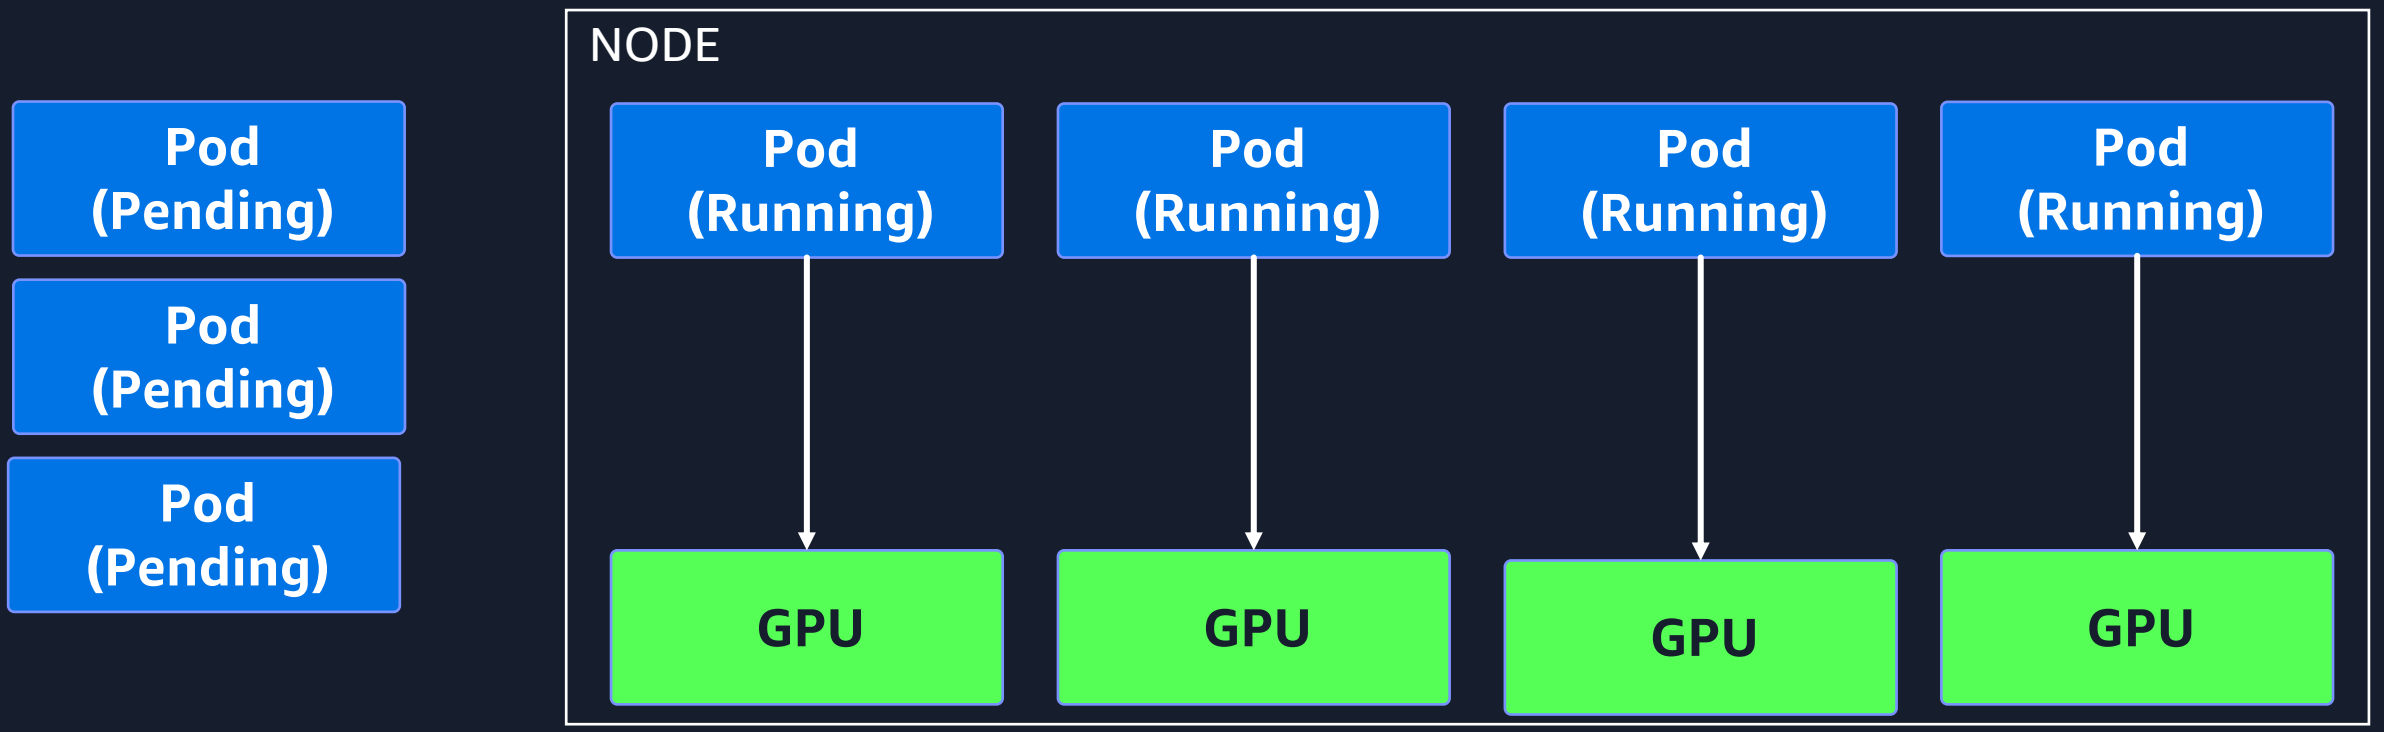

Scheduling Pods

- Allocatable status on Node
- Pods states how much GPU they need
- As Pod wait for GPU, they are pending

```
containers:
- name: cuda-container
  resources:
    limits:
      nvidia.com/gpu: 1 # requesting 1 GPU
```

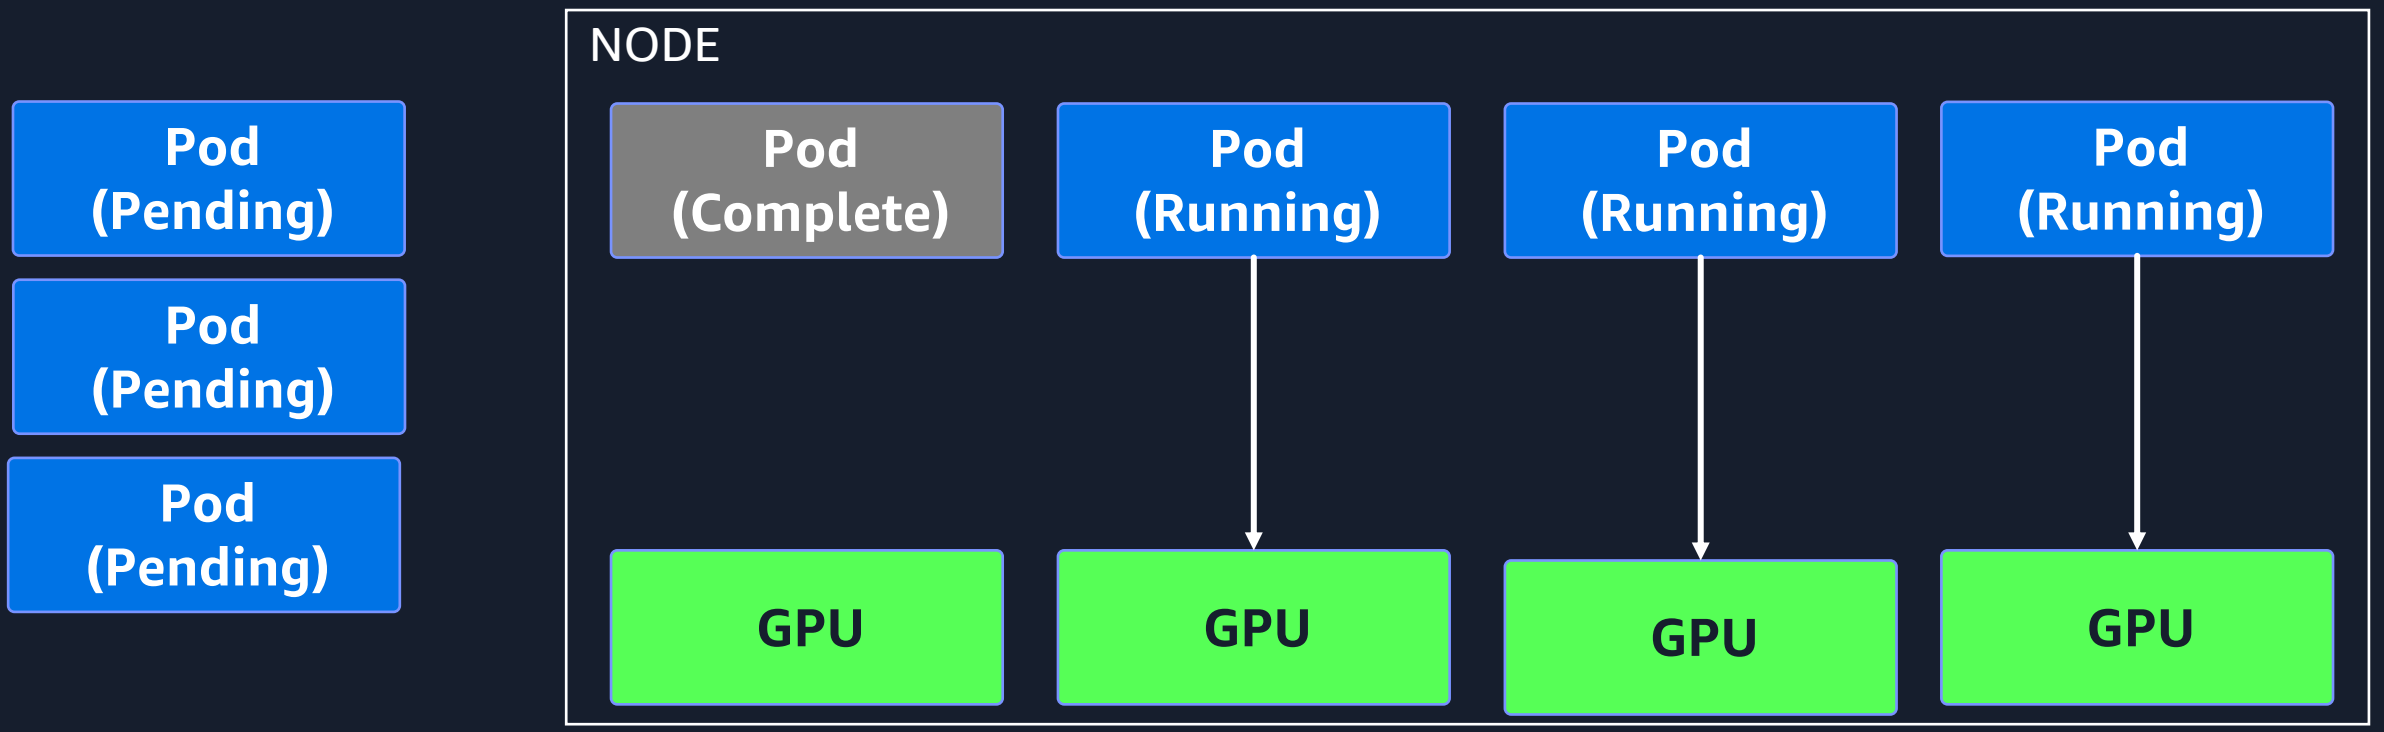

Scheduling Pods

- Allocatable status on Node
- Pods states how much GPU they need
- As Pod wait for GPU, they are pending

```
containers:
- name: cuda-container
  resources:
    limits:
      nvidia.com/gpu: 1 # requesting 1 GPU
```

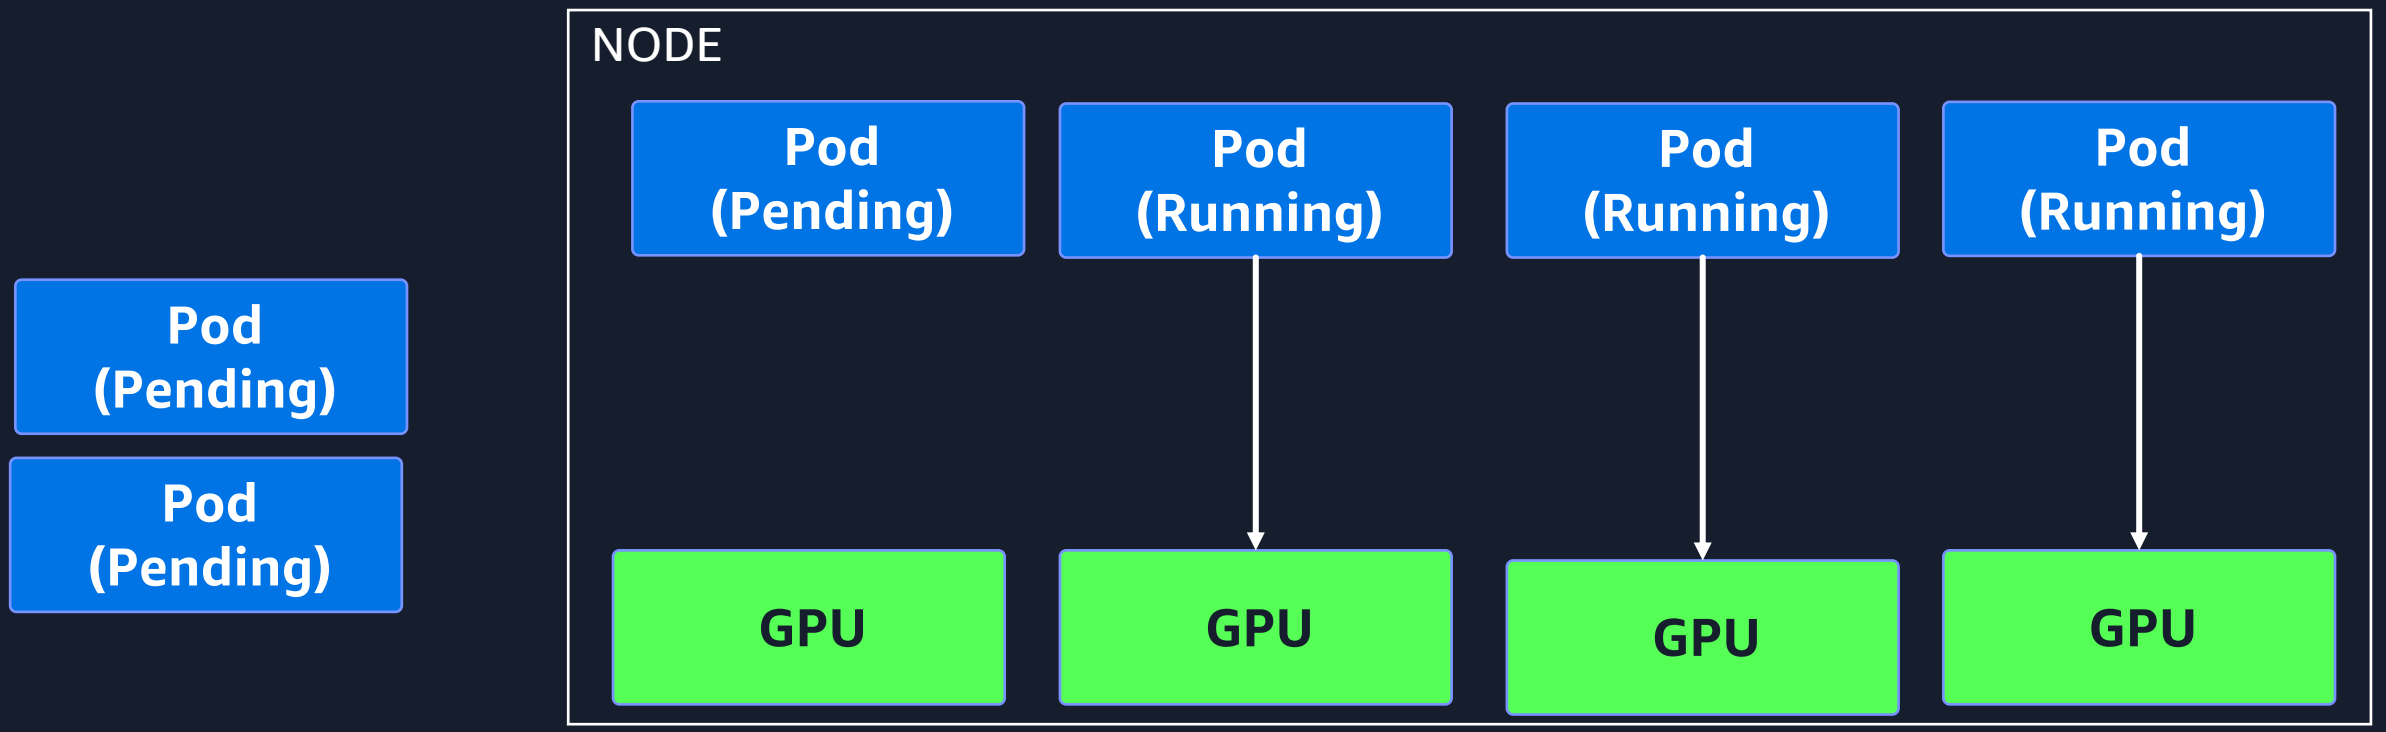

Scheduling Pods

- Allocatable status on Node
- Pods states how much GPU they need
- As Pod wait for GPU, they are pending

```
containers:
- name: cuda-container
  resources:
    limits:
      nvidia.com/gpu: 1 # requesting 1 GPU
```

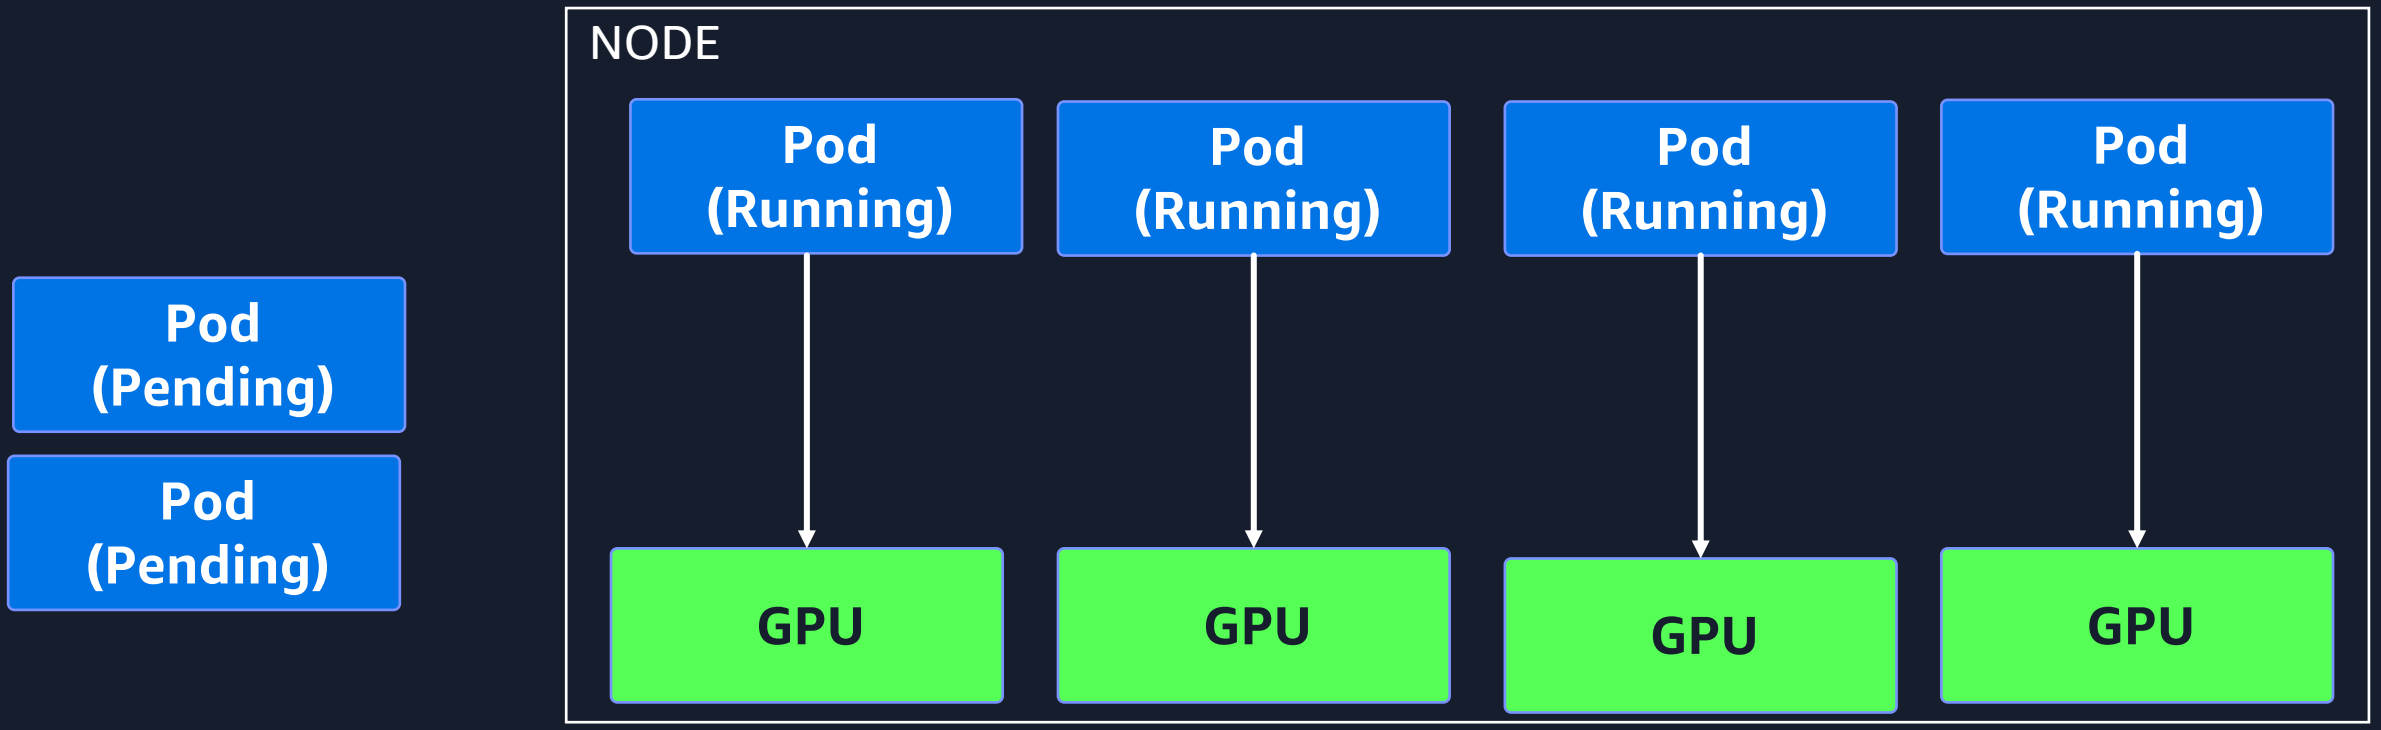

Why and How

Run Containers

- Device Driver
- Container Toolkit
- CUDA

Orchestrate Containers

kubernetes

- Device Plugin
- Node Feature Discovery
- GPU Feature Discovery

Kubernetes Home lab

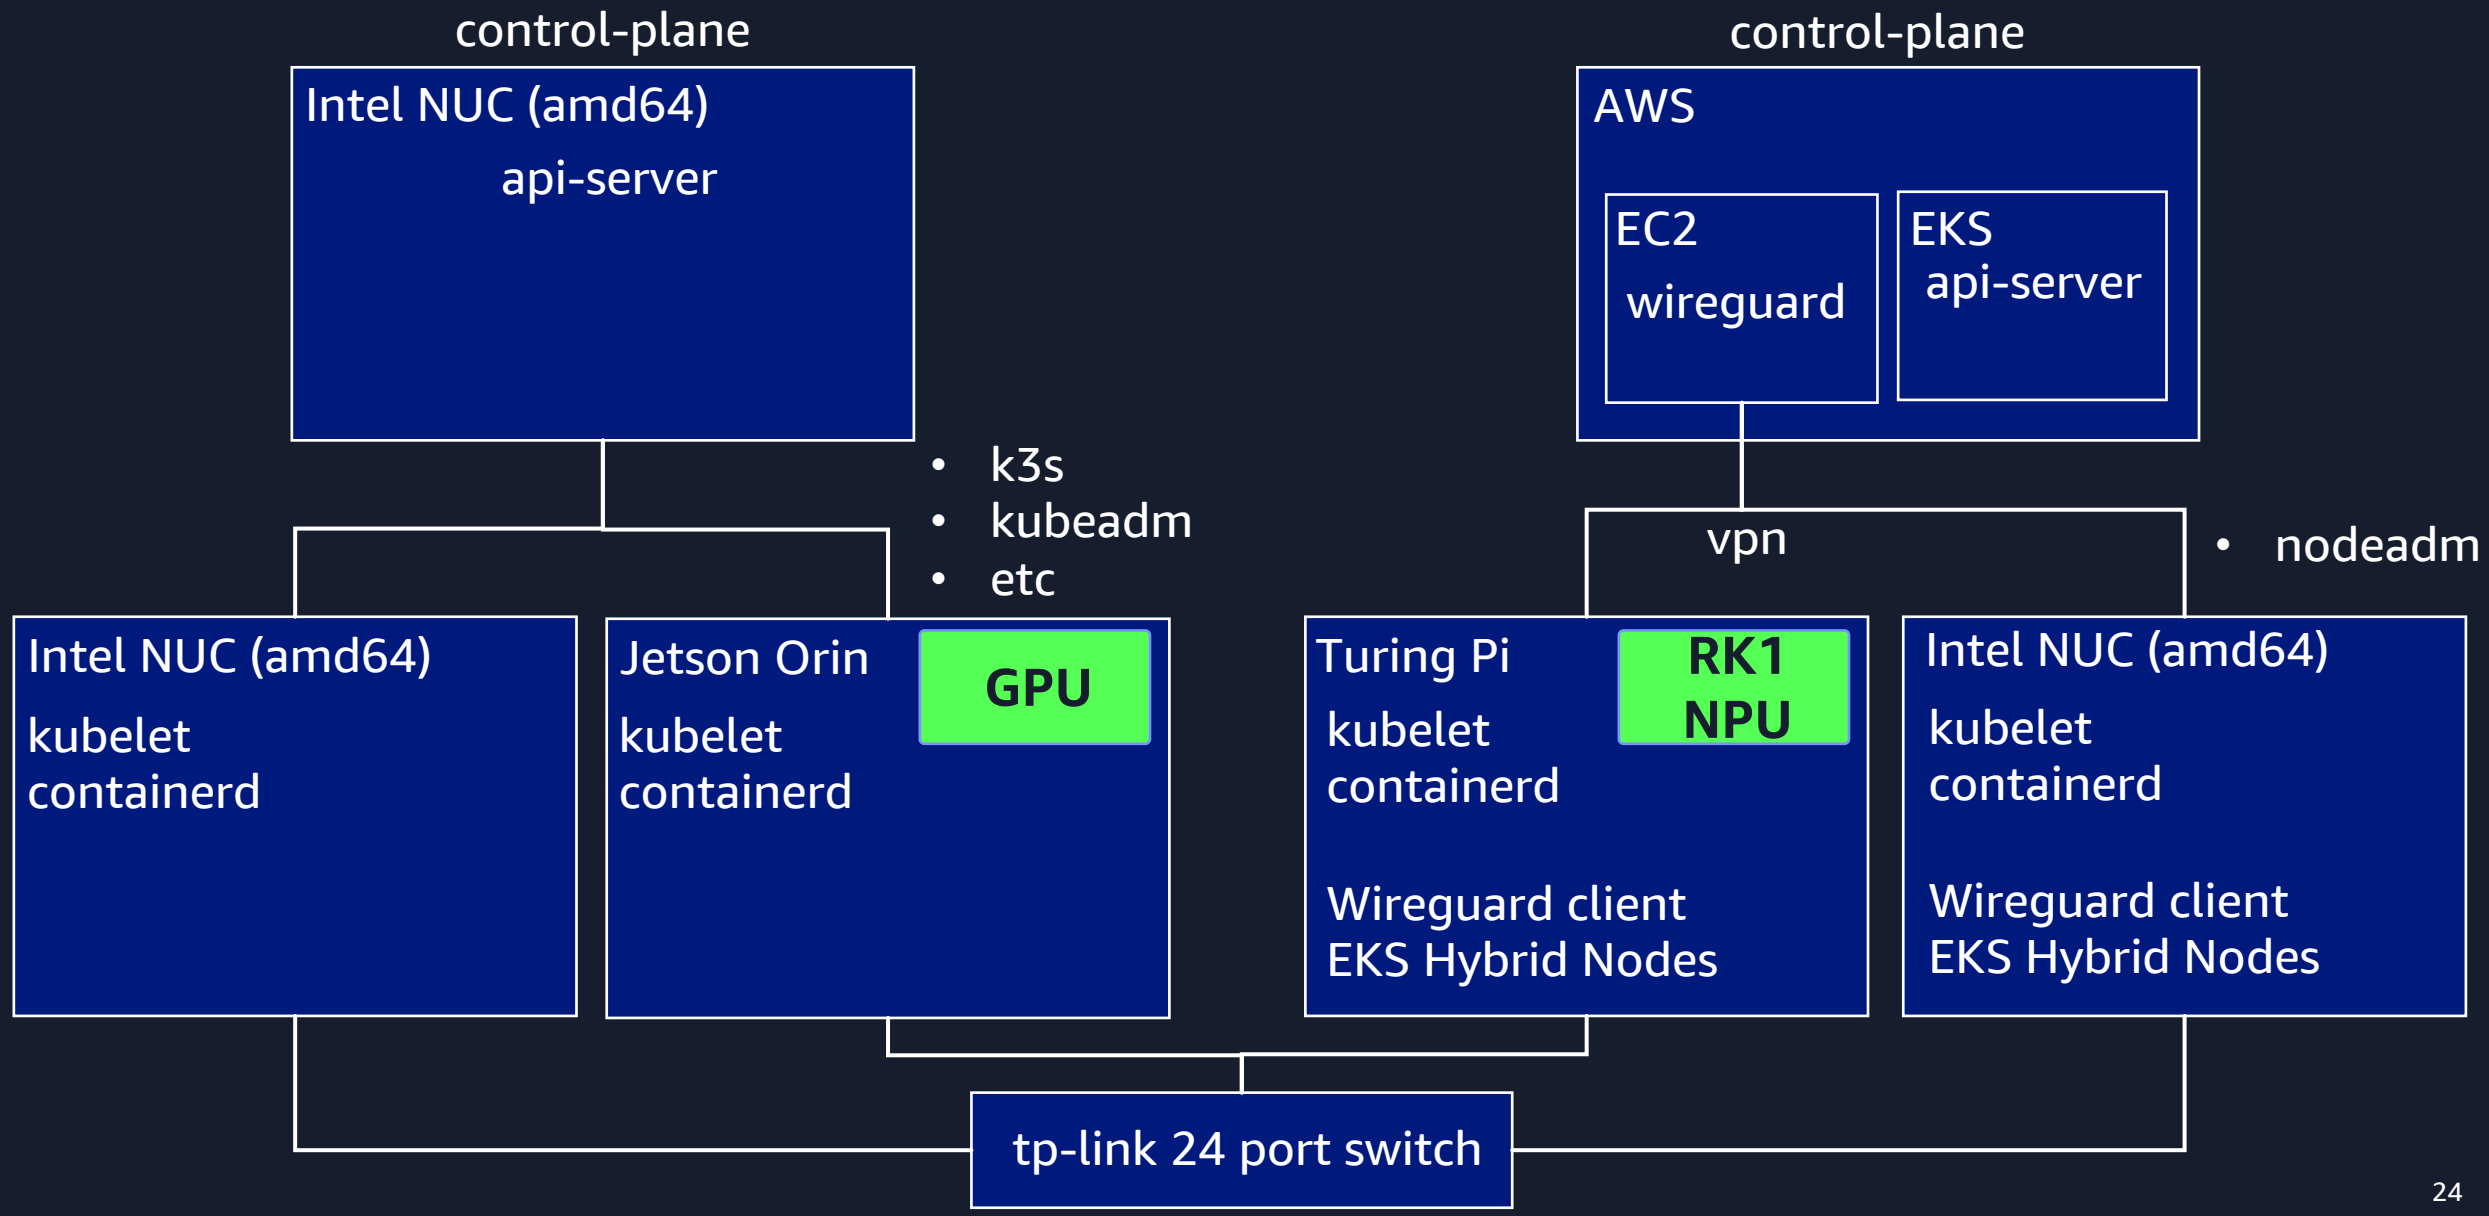

Deep Seek R1 on Jetson kubernetes

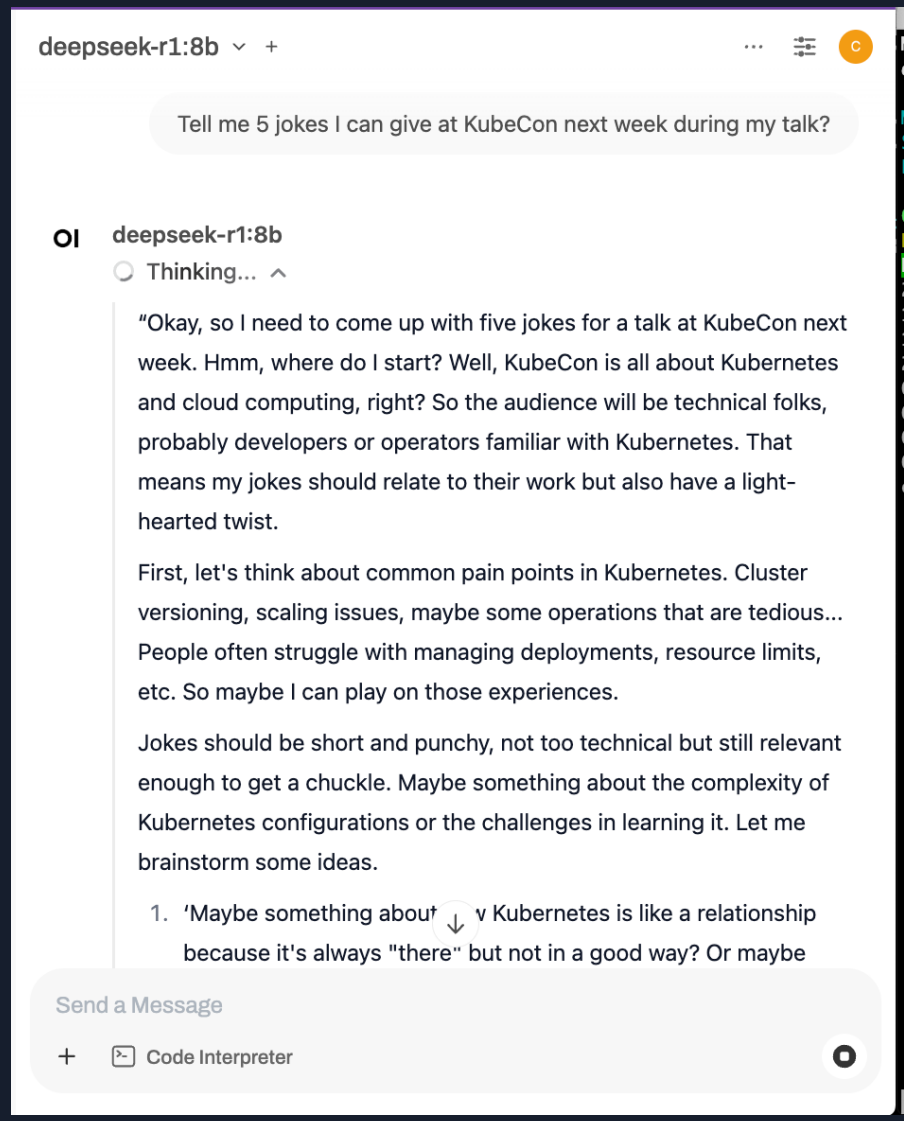

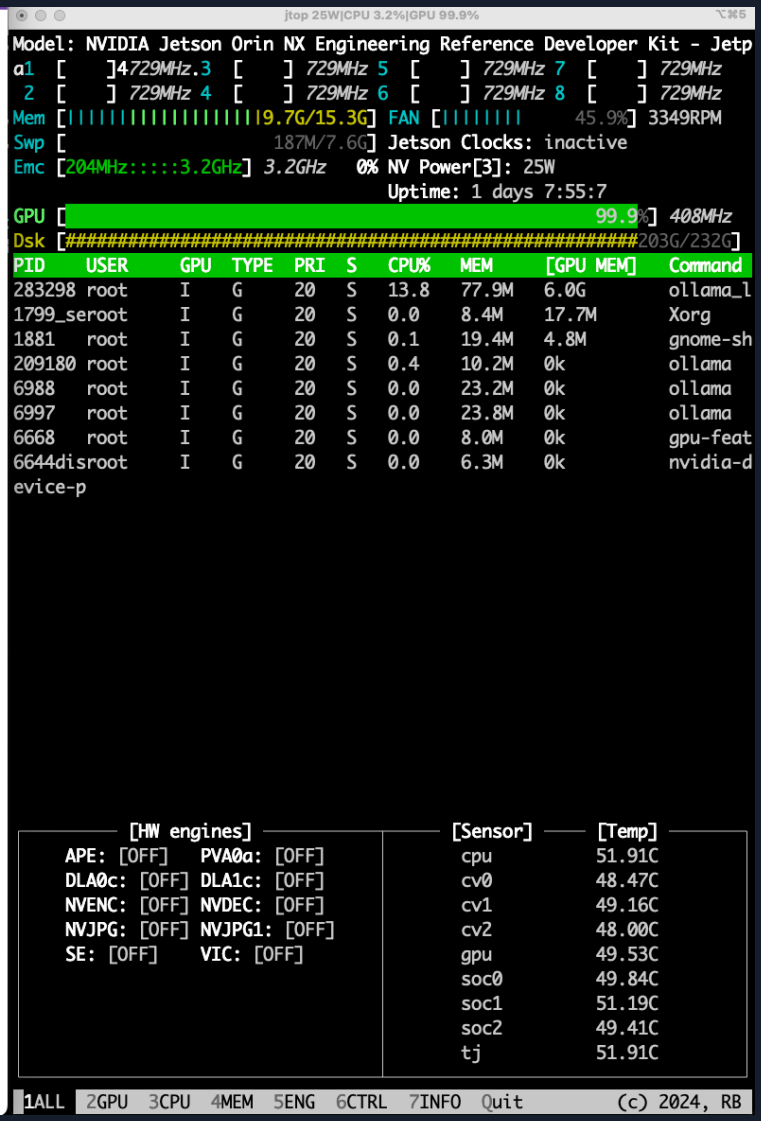

Resources

- https://github.com/csantanapr/k8s-nvidia

AWS Booth
S300

Thank you!

Complete the Survey

Carlos Santana

csantanapr

santana.dev