



## Europe 2025



# Uncharted Waters: Dynamic Resource Allocation for Networking

Lionel Jouin  
Miguel Duarte Barroso



KubeCon



CloudNativeCon

Europe 2025





Lionel Jouin  
Software Engineer @ Ericsson Software  
Technology (EST)  
@LionelJouin



Miguel Duarte Barroso  
Principal Software Engineer @ Red Hat  
@maiqueb



KubeCon



CloudNativeCon

Europe 2025



Problem & Motivation

Multi-network status quo

Intro to DRA

DRA in the Context of Networking

Existing KEPs

Multi-Network  
Scheduling Use Case

Demo





# Problem

Kubernetes is Opinionated

Each Pod Gets One Interface

All Pods are Interconnected

No NAT !

Micro-Segmentation via Network Policies



What if ...

- You don't want the cluster Default Network ?
- You want \*more\* than one Interface ?

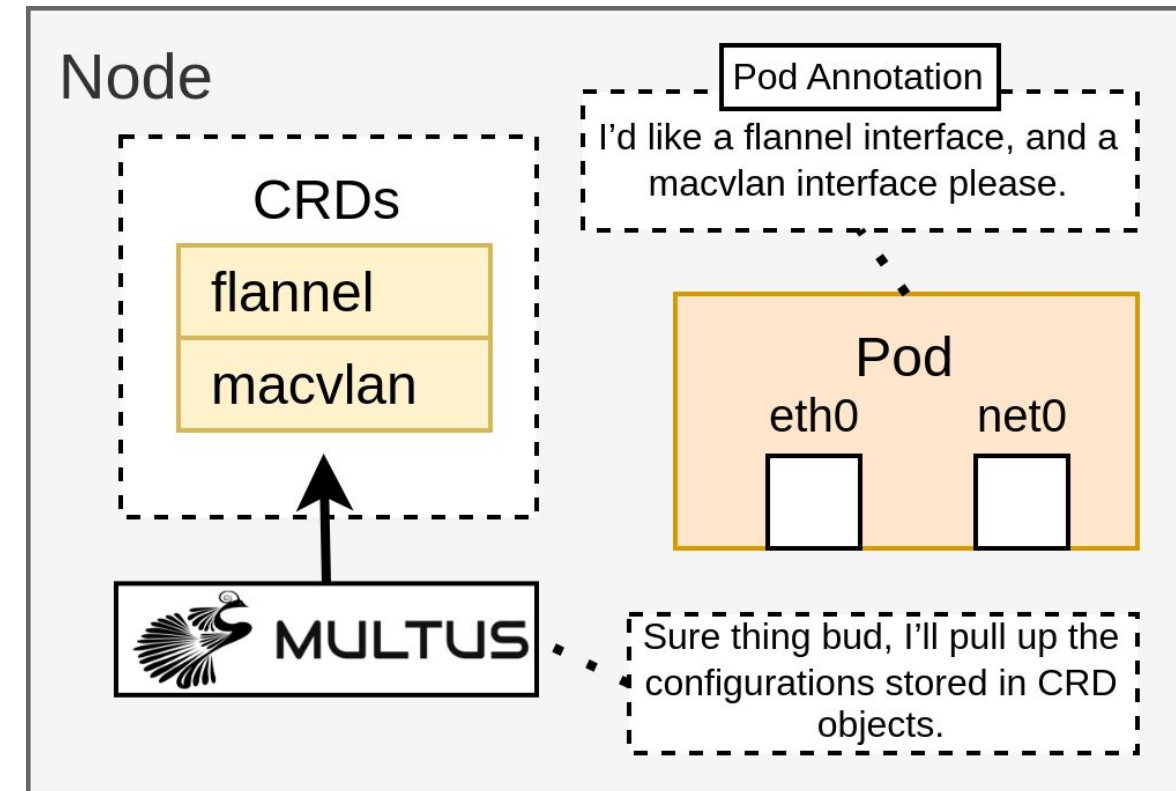


# The Current Status Quo

Network Plumbing Working Group / Multus CNI  
*CNI JSON inside YAML and "out-of-tree"*

```
kind: Pod
metadata:
  name: pod1
  annotations:
    k8s.cni.cncf.io/networks: '{"name":
"media-net", "interface": "net0"}'
spec:
  [...]
```

```
kind: NetworkAttachmentDefinition
[...]
spec:
  conf: '{
    "type": "bridge",
    "name": "media-net",
    "ipam": {"type": "static"}
  }'
```



# KEP-3698: Bringing Multi-Network “In Tree”



KubeCon



CloudNativeCon

Europe 2025

Network Plumbing Working Group / Multus CNI  
*CNI JSON inside YAML and "out-of-tree"*

```
kind: Pod
metadata:
  name: pod1
  annotations:
    k8s.cni.cncf.io/networks: '{"name":
"media-net", "interface": "net0"}'
spec:
  [...]
```

```
kind: NetworkAttachmentDefinition
[...]
spec:
  conf: '{
    "type": "bridge",
    "name": "media-net",
    "ipam": {"type": "static"}
  }'
```


K8s Native Multi-networking / PodNetwork  
*All YAML and all "in tree"*

```
kind: Pod
metadata:
  name: pod1
spec:
  networks:
  - podNetworkName: media-net
    interfaceName: net0
  [...]
```

```
apiVersion: v1
kind: PodNetwork
metadata:
  name: media-net
spec:
  provider:
    "myprovider.io/application"
```



*"Do Not Touch  
the Pod Object  
Specification"*



```
kind: Pod
metadata:
  name: pod1
spec:
  networks:
  - podNetworkName: media-net
    interfaceName: net0
  [...]
```

# The Current Status Quo - Device Plugins

	Device Plugins
Introduced in...	K8s 1.8 (alpha)
Resource Model	Extended Resources: <code>vendor.com/device=X</code>
Sharing across pods?	✗ No (one pod per device)
Per node daemon required?	✓ Yes (device plugin on each node)

# Multi-Network Decomposed




KubeCon



CloudNativeCon

Europe 2025





# What is DRA



KubeCon

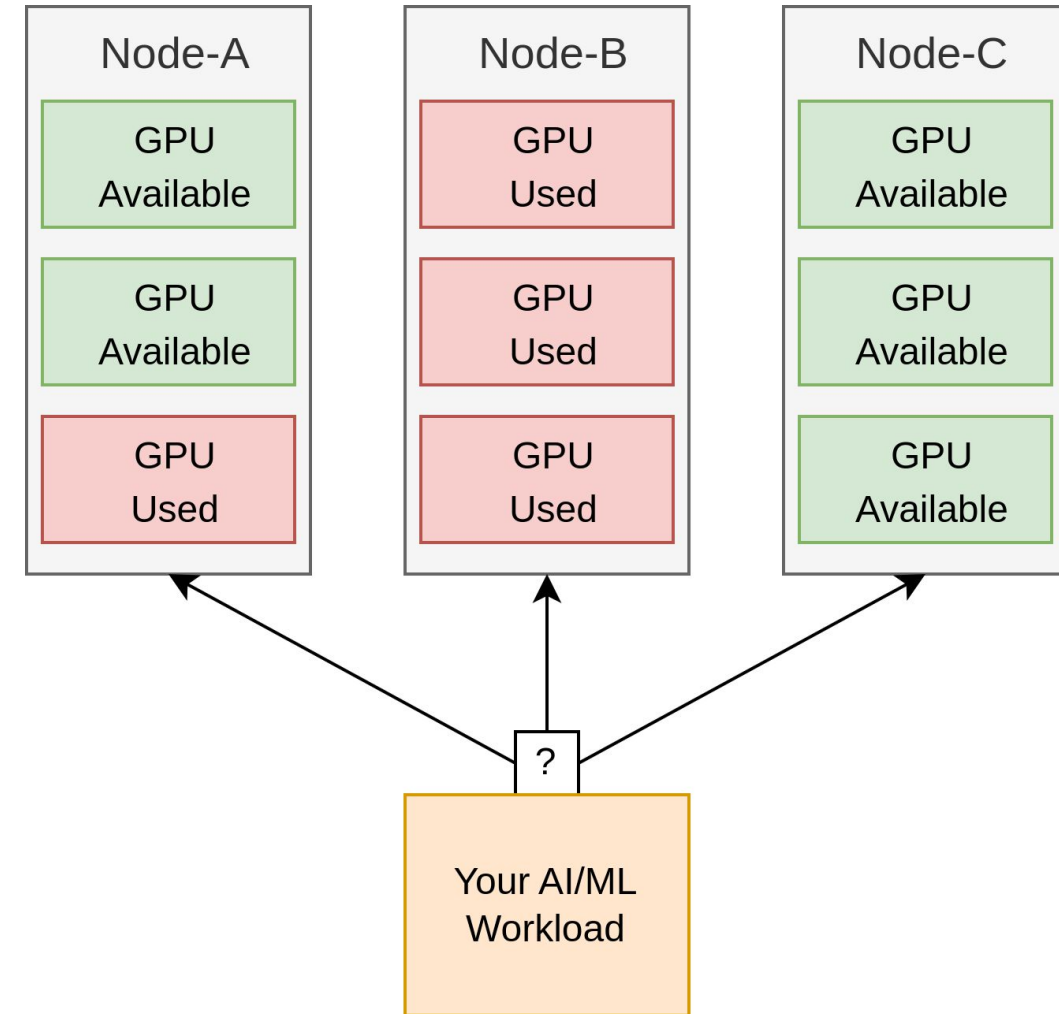


CloudNativeCon

Europe 2025

“an API for requesting and sharing resources between pods and containers inside a pod” - [K8s.io](https://k8s.io)


Currently beta in 1.32





# Generalizing DRA

... an API for requesting and sharing  
***network interfaces*** between pods and  
containers



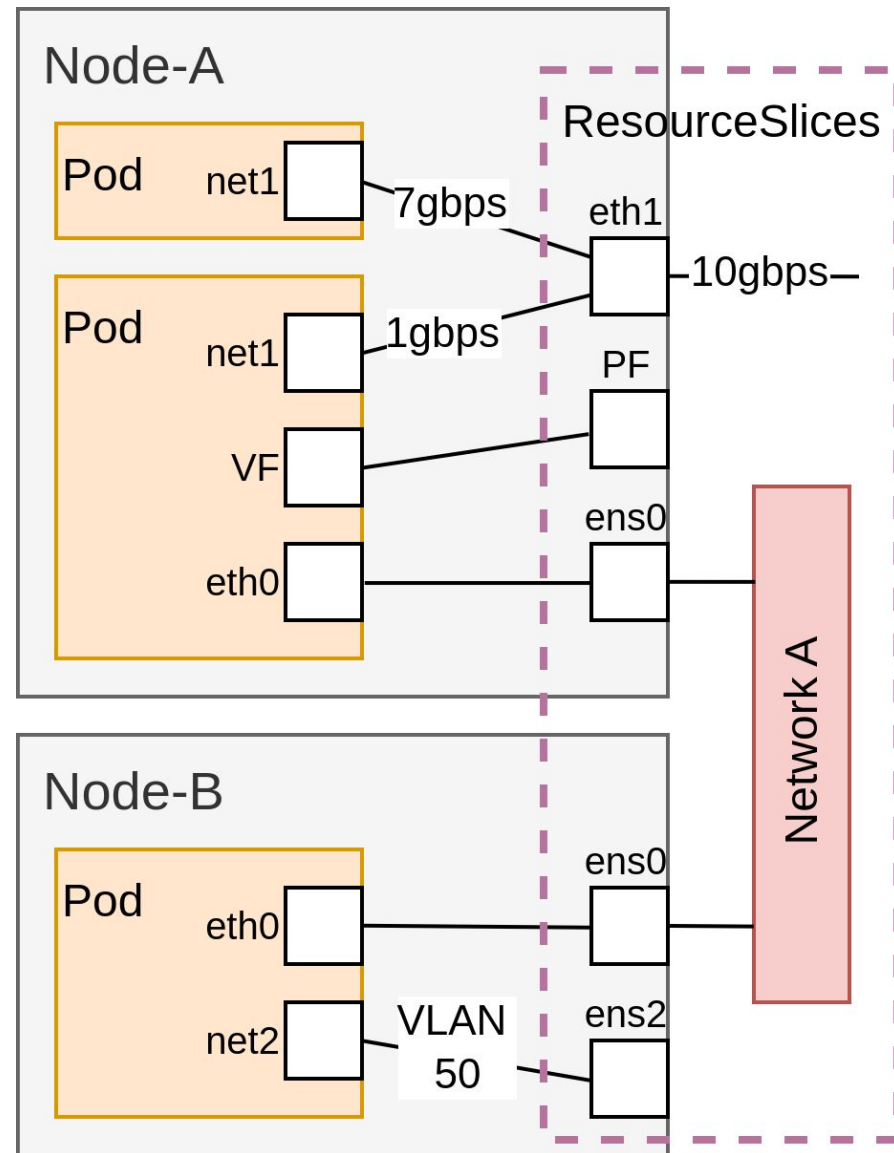
ResourceSlice Lists What is Available

Pod Requests Requirements

Schedule Pods Based on Their Needs

Non-Uniform Clusters

KEP-5075 - Consumable Capacity



# Network Interface Workflow Configuration



KubeCon



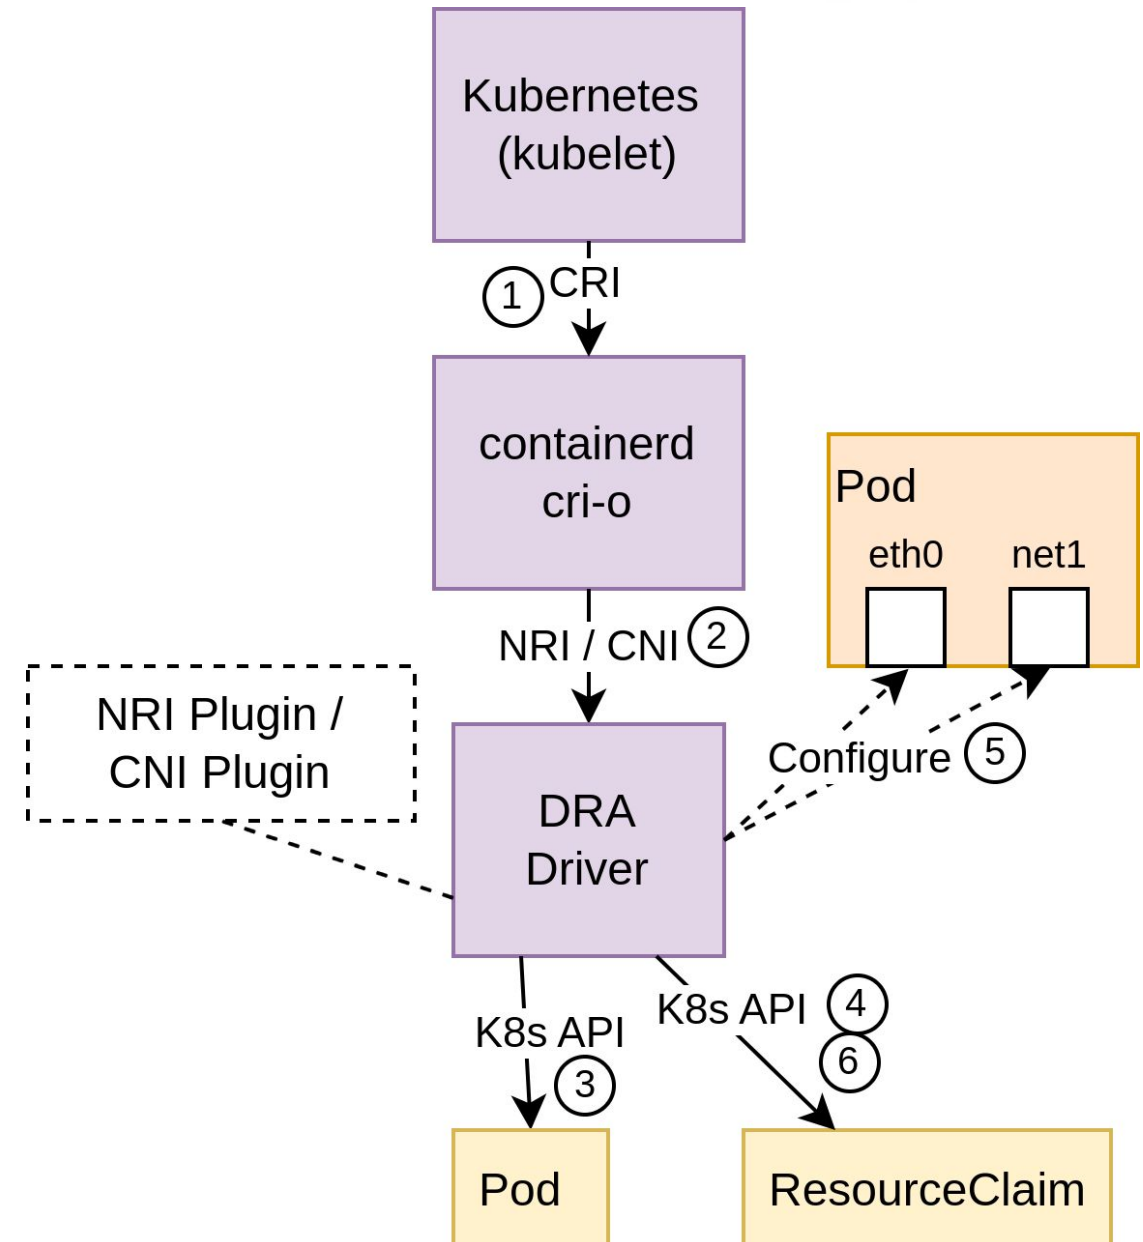
CloudNativeCon

Europe 2025

## Node Resource Interface (NRI)

## Hook into Pod Lifecycle Events

1. RunPodSandbox
2. RunPodSandbox (NRI) / CNI ADD
3. Get Pod
4. Get Pod's ResourceClaims
5. Configure Pod's Network Interfaces
6. Set ResourceClaims Status



## KEP-4817 - ResourceClaim Device Status

### Responsibility on the DRA-Driver

1. Network Services
2. Troubleshooting

### 1 Pod per ResourceClaim

### Alpha in 1.32 and Beta in 1.33

```
apiVersion: resource.k8s.io/v1beta1
kind: ResourceClaim
metadata:
  name: macvlan-eth0
spec:
  ...
status:
  ...
  devices:
  - conditions: []
    device: macvlan-eth0
    driver: cni.dra.driver
    pool: kind-worker
    data:
      cniVersion: 1.0.0
      interfaces:
      - mac: 00:01:ec:f9:12:3b
        name: net1
        sandbox: /var/run/netns/cni-d36910c7
      ips:
      - address: 10.10.1.2/24
        gateway: 10.10.1.1
        interface: 0
    networkData:
      ips:
      - 10.10.1.2/24
      hardwareAddress: 00:01:ec:f9:12:3b
      interfaceName: net1
```



# Demo - MACVLAN on Non-Uniform Cluster


Check the Environment

List Resources Available (ResourceSlices)

Request Resources from Pods (15 replicas)  
- MACVLAN on Top of eth1 Only Available on Kind-Worker

Verify Resources Allocated

Inspect Acquired Resource



# Demo - Environment (Cluster)

```
~$ kubectl get pods --all-namespaces -o wide
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE	IP	NODE
default	cni-dra-driver-57m2k	1/1	Running	0	24s	172.18.0.2	kind-worker2
default	cni-dra-driver-wldbs	1/1	Running	0	24s	172.18.0.4	kind-control-plane
default	cni-dra-driver-zjm9x	1/1	Running	0	24s	172.18.0.3	kind-worker
kube-system	coredns-668d6bf9bc-cb27b	1/1	Running	0	30m	10.244.0.2	kind-control-plane
kube-system	coredns-668d6bf9bc-rsvfn	1/1	Running	0	30m	10.244.0.4	kind-control-plane
kube-system	etcd-kind-control-plane	1/1	Running	0	31m	172.18.0.4	kind-control-plane
kube-system	install-cni-plugins-bghjt	1/1	Running	0	24s	172.18.0.4	kind-control-plane
kube-system	install-cni-plugins-jvrl8	1/1	Running	0	24s	172.18.0.3	kind-worker
kube-system	install-cni-plugins-vstlp	1/1	Running	0	24s	172.18.0.2	kind-worker2
kube-system	kindnet-2k2mj	1/1	Running	0	30m	172.18.0.3	kind-worker
kube-system	kindnet-bbssj	1/1	Running	0	30m	172.18.0.4	kind-control-plane
kube-system	kindnet-f6zx9	1/1	Running	0	30m	172.18.0.2	kind-worker2
kube-system	kube-apiserver-kind-control-plane	1/1	Running	0	31m	172.18.0.4	kind-control-plane
kube-system	kube-controller-manager-kind-control-plane	1/1	Running	0	31m	172.18.0.4	kind-control-plane
kube-system	kube-proxy-hgls4	1/1	Running	0	30m	172.18.0.2	kind-worker2
kube-system	kube-proxy-j49fc	1/1	Running	0	30m	172.18.0.3	kind-worker
kube-system	kube-proxy-lbfhl	1/1	Running	0	30m	172.18.0.4	kind-control-plane
kube-system	kube-scheduler-kind-control-plane	1/1	Running	0	31m	172.18.0.4	kind-control-plane
local-path-storage	local-path-provisioner-58cc7856b6-x774j	1/1	Running	0	30m	10.244.0.3	kind-control-plane

```
~$ docker exec -it kind-worker ip -br link show
```

```
lo            UNKNOWN      00:00:00:00:00:00 <LOOPBACK,UP,LOWER_UP>
eth0@if45     UP           c6:6f:2a:06:54:3f <BROADCAST,MULTICAST,UP,LOWER_UP>
kube-ipvs0    DOWN        46:ad:01:b5:0a:b2 <BROADCAST,NOARP>
eth1@eth0     DOWN        02:a9:d4:b9:b0:17 <BROADCAST,MULTICAST>
```

```
~$ docker exec -it kind-worker2 ip -br link show
```

```
lo            UNKNOWN      00:00:00:00:00:00 <LOOPBACK,UP,LOWER_UP>
eth0@if44     UP           ce:05:77:12:ae:e1 <BROADCAST,MULTICAST,UP,LOWER_UP>
kube-ipvs0    DOWN        96:7f:36:af:46:4b <BROADCAST,NOARP>
```

# Demo - Environment (ResourceSlices)

```
~$ kubectl get resourceslice
```

NAME	NODE	DRIVER	POOL	AGE
kind-worker-cni-dra-driver	kind-worker	cni.dra.networking.x-k8s.io	kind-worker	38s
kind-worker2-cni-dra-driver	kind-worke2	cni.dra.networking.x-k8s.io	kind-worker	38s

```
~$ kubectl get resourceslice
```

```
kind-worker-cni-dra-driver -o yaml
apiVersion: resource.k8s.io/v1beta1
kind: ResourceSlice
metadata:
  name: kind-worker-cni-dra-driver
spec:
```

```
  devices:
  - basic:
      attributes:
        name:
          string: eth0
      name: eth0
  - basic:
      attributes:
        name:
          string: eth1
      name: eth1
```

```
  driver: cni.dra.networking.x-k8s.io
  nodeName: kind-worker
  pool:
    generation: 0
    name: kind-worker
    resourceSliceCount: 1
```

```
~$ kubectl get resourceslice kind-worker2-cni-dra-driver -o
yaml
```

```
apiVersion: resource.k8s.io/v1beta1
kind: ResourceSlice
metadata:
  name: kind-worker2-cni-dra-driver
spec:
```

```
  devices:
  - basic:
      attributes:
        name:
          string: eth0
      name: eth0
```

```
  driver: cni.dra.networking.x-k8s.io
  nodeName: kind-worke2
  pool:
    generation: 0
    name: kind-worker
    resourceSliceCount: 1
```

# Demo (ResourceClaimTemplate + Deployment)



KubeCon



CloudNativeCon

Europe 2025

```
~$ kubectl get resourceclaimtemplate macvlan-eth1-attachment -o yaml
apiVersion: resource.k8s.io/v1beta1
kind: ResourceClaimTemplate
metadata:
  name: macvlan-eth1-attachment
  namespace: default
spec:
  spec:
    devices:
      config:
      - opaque:
          driver: cni.dra.networking.x-k8s.io
          parameters:
            apiVersion: cni.networking.x-k8s.io/v1alpha1
            config:
              cniVersion: 1.0.0
              name: macvlan-eth1
              plugins:
              - ipam:
                  ranges:
                  - - subnet: 10.10.1.0/24
                    type: host-local
                  master: eth1
                  mode: bridge
                  type: macvlan
            ifName: net1
            kind: CNI
          requests:
          - macvlan-eth1
      requests:
      - adminAccess: true
        allocationMode: ExactCount
        count: 1
        deviceClassName: cni.networking.x-k8s.io
        name: macvlan-eth1
      selectors:
      - cel:
          expression: device.driver == "cni.dra.networking.x-k8s.io"
      - cel:
          expression: device.attributes["cni.dra.networking.x-k8s.io"].name == "eth1"
```

```
~$ kubectl apply -f deployment-template.yaml
resourceclaimtemplate.resource.k8s.io/macvlan-eth1-attachment created
deployment.apps/demo-application created
```

```
~$ kubectl get deployment demo-application -o yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  generation: 1
  labels:
    app: demo-application
  name: demo-application
  namespace: default
spec:
  replicas: 15
  selector:
    matchLabels:
      app: demo-application
  template:
    metadata:
      labels:
        app: demo-application
    spec:
      containers:
      - command:
        - sleep
        - infinity
        image: alpine:latest
        imagePullPolicy: IfNotPresent
        name: alpine
      resourceClaims:
      - name: macvlan-eth1-attachment
        resourceClaimTemplateName: macvlan-eth1-attachment
  status:
    availableReplicas: 15
    readyReplicas: 15
    replicas: 15
    updatedReplicas: 15
```



# Demo (ResourceClaims)

```
~$ kubectl get resourceclaim
```

NAME	STATE	AGE
demo-application-c86fd78c5-2lxjq-macvlan-eth1-attachment-kvnjg	allocated, reserved	27s
demo-application-c86fd78c5-65gj7-macvlan-eth1-attachment-q2q7c	allocated, reserved	27s
demo-application-c86fd78c5-672rj-macvlan-eth1-attachment-lmx5g	allocated, reserved	27s
demo-application-c86fd78c5-9msd5-macvlan-eth1-attachment-6xkdq	allocated, reserved	27s
demo-application-c86fd78c5-blzln-macvlan-eth1-attachment-h4rcg	allocated, reserved	27s
demo-application-c86fd78c5-bsndz-macvlan-eth1-attachment-vfq55	allocated, reserved	27s
demo-application-c86fd78c5-d8q26-macvlan-eth1-attachment-btlbh	allocated, reserved	27s
demo-application-c86fd78c5-dwzrt-macvlan-eth1-attachment-5h6nr	allocated, reserved	27s
demo-application-c86fd78c5-fv8sg-macvlan-eth1-attachment-n5f5h	allocated, reserved	27s
demo-application-c86fd78c5-gqs4d-macvlan-eth1-attachment-wcsx9	allocated, reserved	27s
demo-application-c86fd78c5-kbdrm-macvlan-eth1-attachment-fnn5b	allocated, reserved	27s
demo-application-c86fd78c5-mm84g-macvlan-eth1-attachment-4sgx8	allocated, reserved	27s
demo-application-c86fd78c5-nnbxr-macvlan-eth1-attachment-slqxq	allocated, reserved	27s
demo-application-c86fd78c5-qbmwg-macvlan-eth1-attachment-7t9fw	allocated, reserved	27s
demo-application-c86fd78c5-zl248-macvlan-eth1-attachment-4vxzq	allocated, reserved	27s

# Demo (ResourceClaim)

```
~$ kubectl get resourceclaim -o yaml demo-application-c86fd78c5-2lxjq-macvlan-eth1-attachment-kvnjg
```

```
apiVersion: resource.k8s.io/v1beta1
kind: ResourceClaim
metadata:
  annotations:
    resource.kubernetes.io/pod-claim-name: macvlan-eth1-attachment
  generateName: demo-application-c86fd78c5-2lxjq-macvlan-eth1-attachment-
  name: demo-application-c86fd78c5-2lxjq-macvlan-eth1-attachment-kvnjg
  namespace: default
  ownerReferences:
  - apiVersion: v1
    blockOwnerDeletion: true
    controller: true
    kind: Pod
    name: demo-application-c86fd78c5-2lxjq
    uid: a2735b69-92de-4ac1-8461-725323ddefb7
  resourceVersion: "3618"
  uid: 80067ad8-ba85-4230-bc4a-ee3e1e1c4a2b
[...]
```

```
[...]
spec:
  devices:
    config:
      - opaque:
          driver: cni.dra.networking.x-k8s.io
          parameters:
            apiVersion: cni.networking.x-k8s.io/v1alpha1
            config:
              cniVersion: 1.0.0
              name: macvlan-eth1
              plugins:
                - ipam:
                    ranges:
                      - - subnet: 10.10.1.0/24
                        type: host-local
                    master: eth1
                    mode: bridge
                    type: macvlan
              ifName: net1
              kind: CNI
          requests:
            - macvlan-eth1
  requests:
  - adminAccess: true
    allocationMode: ExactCount
    count: 1
    deviceClassName: cni.networking.x-k8s.io
    name: macvlan-eth1
    selectors:
      - cel:
          expression: device.driver == "cni.dra.networking.x-k8s.io"
      - cel:
          expression: device.attributes["cni.dra.networking.x-k8s.io"].name == "eth1"
[...]
```

# Demo (ResourceClaim Status)

```
[...]
status:
  allocation:
    devices:
      config:
        - opaque:
            driver: cni.dra.networking.x-k8s.io
            parameters:
              apiVersion: cni.networking.x-k8s.io/v1alpha1
              config:
                cniVersion: 1.0.0
                name: macvlan-eth1
                plugins:
                  - ipam:
                      ranges:
                        - - subnet: 10.10.1.0/24
                          type: host-local
                      master: eth1
                      mode: bridge
                      type: macvlan
            ifName: net1
            kind: CNI
          requests:
            - macvlan-eth1
            source: FromClaim
          results:
            - adminAccess: true
              device: eth1
              driver: cni.dra.networking.x-k8s.io
              pool: kind-worker
              request: macvlan-eth1
      nodeSelector:
        nodeSelectorTerms:
          - matchFields:
              - key: metadata.name
                operator: In
                values:
                  - kind-worker
```

[...]

```
[...]
devices:
  - conditions: null
    data:
      cniVersion: 1.0.0
      interfaces:
        - mac: 8e:18:6b:cb:d3:c5
          name: net1
          sandbox: /var/run/netns/cni-ed50d613-a4a9-1b5b-c422-3e620dc17878
      ips:
        - address: 10.10.1.16/24
          gateway: 10.10.1.1
          interface: 0
      device: eth1
      driver: cni.dra.networking.x-k8s.io
      networkData:
        hardwareAddress: 8e:18:6b:cb:d3:c5
        interfaceName: net1
        ips:
          - 10.10.1.16/24
      pool: kind-worker
    reservedFor:
      - name: demo-application-c86fd78c5-21xjq
        resource: pods
        uid: a2735b69-92de-4ac1-8461-725323ddefb7
```

# Demo (Pods + Network Interfaces)

```
~$ kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
cni-dra-driver-57m2k	1/1	Running	0	105s	172.18.0.2	kind-worker2	<none>	<none>
cni-dra-driver-wldbs	1/1	Running	0	105s	172.18.0.4	kind-control-plane	<none>	<none>
cni-dra-driver-zjm9x	1/1	Running	0	105s	172.18.0.3	kind-worker	<none>	<none>
demo-application-c86fd78c5-2lxjq	1/1	Running	0	47s	10.244.2.16	kind-worker	<none>	<none>
demo-application-c86fd78c5-65gj7	1/1	Running	0	47s	10.244.2.9	kind-worker	<none>	<none>
demo-application-c86fd78c5-672rj	1/1	Running	0	47s	10.244.2.8	kind-worker	<none>	<none>
demo-application-c86fd78c5-9msd5	1/1	Running	0	47s	10.244.2.10	kind-worker	<none>	<none>
demo-application-c86fd78c5-blzln	1/1	Running	0	47s	10.244.2.11	kind-worker	<none>	<none>
demo-application-c86fd78c5-bsndz	1/1	Running	0	47s	10.244.2.4	kind-worker	<none>	<none>
demo-application-c86fd78c5-d8q26	1/1	Running	0	47s	10.244.2.14	kind-worker	<none>	<none>
demo-application-c86fd78c5-dwzrt	1/1	Running	0	47s	10.244.2.15	kind-worker	<none>	<none>
demo-application-c86fd78c5-fv8sg	1/1	Running	0	47s	10.244.2.13	kind-worker	<none>	<none>
demo-application-c86fd78c5-gqs4d	1/1	Running	0	47s	10.244.2.12	kind-worker	<none>	<none>
demo-application-c86fd78c5-kbdrm	1/1	Running	0	47s	10.244.2.2	kind-worker	<none>	<none>
demo-application-c86fd78c5-mm84g	1/1	Running	0	47s	10.244.2.6	kind-worker	<none>	<none>
demo-application-c86fd78c5-nnbxr	1/1	Running	0	47s	10.244.2.5	kind-worker	<none>	<none>
demo-application-c86fd78c5-qbmwg	1/1	Running	0	47s	10.244.2.7	kind-worker	<none>	<none>
demo-application-c86fd78c5-zl248	1/1	Running	0	47s	10.244.2.3	kind-worker	<none>	<none>

```
~$ kubectl exec -it demo-application-c86fd78c5-2lxjq -- ip a
```

```
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: eth0@if19: <BROADCAST,MULTICAST,UP,LOWER_UP,M-DOWN> mtu 1500 qdisc noqueue state UP
    link/ether 9a:a4:24:47:f2:01 brd ff:ff:ff:ff:ff:ff
    inet 10.244.2.16/24 brd 10.244.2.255 scope global eth0
        valid_lft forever preferred_lft forever
    inet6 fe80::98a4:24ff:fe47:f201/64 scope link
        valid_lft forever preferred_lft forever
3: net1@eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP
    link/ether 8e:18:6b:cb:d3:c5 brd ff:ff:ff:ff:ff:ff
    inet 10.10.1.16/24 brd 10.10.1.255 scope global net1
        valid_lft forever preferred_lft forever
    inet6 fe80::8c18:6bff:feeb:d3c5/64 scope link
        valid_lft forever preferred_lft forever
```



## CNI

- Weekly Meeting (Monday)
- <https://github.com/containernetworking/meeting-notes>

## WG Device Management

- Every Two Weeks (Tuesday)
- <https://docs.google.com/document/d/1qxl87VqGtgN7EAJlqVfxx86HGKEAc2A3SKru8nJHNkQ>

## Multi-Network

- Every Week (Wednesday)
- <https://github.com/kubernetes-sigs/multi-network>

## K8S Network Plumbing WG

- Every Two Weeks (Thursday)
- <https://github.com/k8snetworkplumbingwg>

- DRA Allows Users to Request Arbitrary Resources for Pods
- Network Interfaces Can Be Considered as a Resource
- DRA's Upcoming "[Consumable Capacity](#)" Enables Unique New Use Cases
  - Scheduling / Resource Sharing



KubeCon



CloudNativeCon

Europe 2025

Thank You!